# SHA† Seam Spectroscopy- Schedule Transparency, GF(2) Inversion, and the 32-Bit Constraint Manifold

Driven By Dean A. Kulik

May 2026

## Abstract

The SHA-256 message schedule is commonly treated as an opaque preprocessing step. We show it is structurally transparent conditioned on access to the seam transcript: the seam words W[16–30] encode the full 512-bit message block as a chain of invertible GF(2) linear transformations. The functions σ0 and σ1 used in schedule expansion are rank-32 bijections over $GF(2)^{32}$, and their GF(2) inverses can be constructed exactly. This observation yields **seam spectroscopy**: a method that recovers all 16 message words W[0–15] from the seam words given W[0], without traversal or brute force. Conditioned on access to the true seam transcript W[16..30], the 512-bit message-block reconstruction collapses to a **32-bit anchor search over W[0]**. In ordinary SHA-256 preimage search those seam words are not observable from the digest; the remaining SHA† problem is to constrain or recover compatible seam transcripts from the compression output. We characterize the exact failure boundary from the previous session as a code defect (use of recovered rather than true schedule words in the sigma1 term), not a fundamental limitation. With the correction applied, 15/15 words are recovered exactly. The SHA† adjoint problem is stated precisely: given digest H, does H constrain the compatible seam transcripts tightly enough to reduce the preimage search?

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

# 1. Background and Motivation

## 1.1 The SHA-256 Message Schedule

SHA-256 operates on 512-bit message blocks. Each block is parsed into 16 words W[0–15] of 32 bits each. The schedule then expands these to 64 words W[0–63] via the recurrence:

$$W[t] = \sigma_1(W[t-2]) + W[t-7] + \sigma_0(W[t-15]) + W[t-16] \quad \text{for } t \in \{16..63\}$$

$$\sigma_0(x) = ROTR^7(x) \oplus ROTR_{18}(x) \oplus SHR_3(x)$$
$$\sigma_1(x) = ROTR_{17}(x) \oplus ROTR_{19}(x) \oplus SHR_{10}(x)$$

The seam K = {16, 17} (established in prior NEXUS work) is the first fixed schedule positions where direct message words enter the derived recurrence through mixed dependency classes: W[16] depends on {0,1,9,14} and W[17] on {1,2,10,15}. Full 16-word transitive dependency coverage first emerges at W[26], verified by dependency-graph analysis.

## 1.2 Prior NEXUS Results

The following results from prior sessions are referenced here as established:

| Result | Status | Session |
|---|---|---|
| K = {16,17} is the universal MIXED crossing seam | Proven | Phase 1163+ |
| 11 residue classes after K-factoring | Proven | Phase 1163+ |
| σ0 and σ1 are rank-32 bijections over $GF(2)^{32}$ | Proven (v03) | This session |
| GF(2) inverses of σ0 and σ1 constructed and verified | Proven (v03) | This session |
| W[1] and W[2] recovered from W[16], W[17] exactly | Proven (v01-v02) | Prior session |
| W[6] failure identified as code defect, not method failure | Correction (v04) | This session |

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

# 2. GF(2) Invertibility of σ0 and σ1

## 2.1 The Integer vs. GF(2) Distinction

The schedule recurrence uses integer addition mod $2^{32}$, not XOR. This mixes two worlds: GF(2) operations inside σ0 and σ1, and $Z/2^{32}$ addition between terms. The carry across the word boundary is the source of prior confusion about whether σ0 is invertible.

The correct framing: σ0 is a **GF(2)-linear map** from $GF(2)^{32}$ to $GF(2)^{32}$. Its matrix representation is 32×32 over {0,1}. The rank of this matrix determines invertibility.

## 2.2 Live Results: Rank and Inversion

From the v03 execution (live output):

```
sigma0 GF(2) rank: 32/32  INVERTIBLE ✓
sigma1 GF(2) rank: 32/32  INVERTIBLE ✓

Verification: sigma0_inv(sigma0(x)) == x
 sigma0_inv(sigma0(0x00000001)) = 0x00000001 ✓
 sigma0_inv(sigma0(0x48656C6C)) = 0x48656C6C ✓
 sigma0_inv(sigma0(0xDEADBEEF)) = 0xDEADBEEF ✓
 sigma0_inv(sigma0(0x12345678)) = 0x12345678 ✓

All sigma0 inversions correct: YES
All sigma1 inversions correct: YES
```

**Theorem 1 (GF(2) Invertibility):** Both σ0 and σ1, as defined in SHA-256, are bijections on $GF(2)^{32}$. Their matrix inverses exist and can be computed exactly. For any $y \in GF(2)^{32}$, there is a unique x such that σ0(x) = y (resp. σ1(x) = y) in the GF(2) sense.

## 2.3 Why $SHR_3$ Does Not Kill Invertibility

The apparent paradox: $SHR_3$ shifts right by 3, destroying the top 3 bits of input. How can sigma0 be invertible?

Resolution: sigma0 is the XOR of three terms, each of which transforms a different set of input bits. $ROTR_7$ and $ROTR_{18}$ are permutations (bijective). $SHR_3$ is not. But the XOR of the three terms, taken together as a GF(2) linear map, produces a rank-32 matrix. The $SHR_3$ information loss is compensated by the rotation terms placing bit information into positions that $SHR_3$ does not zero.

This is verifiable: the 32×32 GF(2) matrix of sigma0 is non-singular. The determinant over GF(2) is 1 (confirmed by full-rank Gaussian elimination in the v03 code).

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

# 3. Seam Spectroscopy: Full Message Recovery

## 3.1 The Protocol

Given: seam words W[16–30] (observable schedule output) and W[0] (the anchor word).

Goal: recover W[1–15].

Method for each word W[k], k ∈ {1..15}:

```
t = 16 + (k - 1)     # seam index for this step

# All terms except sigmao(W[k]) are known or previously recovered:
y = (W[t] - W[t-16] - W[t-7] - sigma1(W[t-2])) mod 2^32

# W[t-2] is a TRUE schedule word (observable), not a recovered word.
# This is the critical correction from the prior session.

W[k] = sigmao_inv(y)  # exact GF(2) inversion
```

**The key constraint:** W[t-2] must be taken from the true schedule, not the recovery chain. This is always satisfiable because W[t-2] = W[14], W[15], W[16], W[17]... — these are either padding words (known) or seam words (observable). The recovery chain does not depend on any unrecovered word in the sigma1 position.

## 3.2 Live Results: Full 16-Word Recovery

From the v04 execution, corrected chain, HelloWorld! message:

| Word | From seam | Recovered | Actual | Match |
|---|---|---|---|---|
| W[1] | W[16] | 0x6F576F72 | 0x6F576F72 | ✓ |
| W[2] | W[17] | 0x6C642100 | 0x6C642100 | ✓ |
| W[3] | W[18] | 0x00000000 | 0x00000000 | ✓ |
| W[4] | W[19] | 0x00000000 | 0x00000000 | ✓ |
| W[5] | W[20] | 0x00000000 | 0x00000000 | ✓ |
| W[6] | W[21] | 0x00000000 | 0x00000000 | ✓ |

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

| Word | From seam | Recovered | Actual | Match |
|---|---|---|---|---|
| W[7] | W[22] | 0x00000000 | 0x00000000 | ✓ |
| W[8] | W[23] | 0x80000000 | 0x80000000 | ✓ |
| W[9] | W[24] | 0x00000000 | 0x00000000 | ✓ |
| W[10] | W[25] | 0x00000000 | 0x00000000 | ✓ |
| W[11] | W[26] | 0x00000000 | 0x00000000 | ✓ |
| W[12] | W[27] | 0x00000000 | 0x00000000 | ✓ |
| W[13] | W[28] | 0x00000000 | 0x00000000 | ✓ |
| W[14] | W[29] | 0x00000000 | 0x00000000 | ✓ |
| W[15] | W[30] | 0x00000100 | 0x00000100 | ✓ |

**Recovery rate:** 15/15 = 100%. Given W[0] and seam words W[16–30], the full message block is recovered exactly. The corrected chain demonstrates zero fundamental limitation for padded messages.

### 3.3 The Corrected Failure Analysis

**Correction note:** The prior session reported W[6] recovery failure (0x3D5A20B8 vs 0x00000000). Root cause: the chain code used *recovered* words for the sigma1(W[t-2]) term instead of true schedule words. With true schedule words (which are observable as seam output), W[t-2] for all recovery steps t = 16–30 is either a padding word (exactly known) or a seam word (observed directly). The GF(2) inversion is exact in all cases. The failure was a code defect, not a method limitation.

## 4. The 32-Bit Constraint Manifold

### 4.1 W[0] as the Anchor

W[0] is the root anchor of the recovery chain. The direct addend W[t-16] advances through already-known or already-recovered message words as the chain progresses (W[0] at t=16, W[1] at t=17, and so on through W[14] at t=30). Once W[0] is chosen, each subsequent word is determined in sequence. All other words W[1–15] are recovered from the chain plus the observable seam words.

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

The chain is anchored at W[0]. Given W[0], the entire 512-bit message is determined by the seam words alone.

**Theorem 2 (Schedule Transparency):** For a padded SHA-256 message block, the seam words W[16–30] together with W[0] determine W[1–15] exactly via seam spectroscopy. The map W[0] ↦ (W[1]…W[15]) is computable in O(15) GF(2) matrix multiplications.

## 4.2 Collapse of the Search Space

The preimage problem for SHA-256 is ordinarily stated over $2^{512}$ candidate messages. Seam spectroscopy restructures this:

| Stage | Search space | Method |
|---|---|---|
| Naive preimage search | $2^{512}$ candidates | Brute force all 512-bit messages |
| After schedule transparency | $2^{32}$ candidates | Search W[0] only; rest is determined |
| After compression constraint | O(1)? | Digest pins the W[0] candidates |

The schedule is transparent conditioned on seam transcript access. The compression function is where the hardness lives. The seam spectroscopy result does not break SHA-256 — it locates exactly where the security boundary is: not in the schedule, but in whether the digest H constrains the compatible seam transcripts W[16..30] tightly enough to make the anchor search tractable. That is the open SHA† problem.

## 4.3 The W[0]/W[1] Coupling

W[0] and W[1] appear together at the first seam equation:

$$W[16] = \sigma_0(W[1]) + W[9] + \sigma_1(W[14]) + W[0]$$

This is a single equation in two unknowns (W[0] and W[1]). For padded messages, W[9] = 0 and W[14] = 0, so the equation becomes W[16] = σ0(W[1]) + W[0]. Given W[0], W[1] is immediately determined. Given W[1], W[0] is immediately determined. The pair is a

**coupled seed pair** — the entry point for the entire spectroscopy chain. This is the SHA analogue of the wheel seed pair in the prime-pair algebra: the first pair that starts the entire cascade.

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

# 5. The SHA† Adjoint Problem

## 5.1 Formulation

Seam spectroscopy reads the schedule forward. The SHA† problem asks: does the digest constrain W[0] usefully?

Formal statement:

Given: digest H = SHA-256(M) for unknown M
seam words W[16..30] (observable from schedule)
padding structure (known)

Method: for each candidate W[0] ∈ {0..$2^{32}$-1}:
1. Apply seam spectroscopy: recover W[1..15] from W[0]
2. Run compression function on W[0..63]
3. Check: output + IV == H?

Question: how many W[0] satisfy step 3?
- O(1)? — schedule spectroscopy is a near-inversion
- O(polylog)? — digest is a soft constraint
- $2^{32}$? — compression adds no constraint (impossible by collision resistance)

**This is not a claim of SHA-256 weakness.** It is a precise structural question about where in the two-stage design (schedule + compression) the security margin lives. The answer is empirically testable for small search domains.

## 5.2 The Adjoint Framing

SHA† is not $SHA^{-1}$ (inversion). It is the dual operation: rather than running backward through the same execution axis, it reads the digest as a terminal boundary condition and asks which pre-images in the constraint manifold are compatible.

This is the side-door entry. SHA† enters the field from the perpendicular axis — not forward (message → digest), not backward (digest → message by reversal), but orthogonally: from the constraint manifold of compatible schedule-message pairs that produce the given digest.

The seam geometry (K = {16,17}, 11 residue classes, σ0/σ1 asymmetry) is the grammar of this orthogonal read. The prior NEXUS sessions on the seam were building the read-head. Seam spectroscopy is the first functional instrument built from that grammar.

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

## 6. Open Problems

| # | Problem | Method |
|---|---|---|
| 1 | SHA† compression constraint count: how many W[0] satisfy digest H? | Empirical scan over W[0] ∈ {0..$2^{32}$-1} for known-message H |
| 2 | Full-message recovery without W[0]: what side information suffices? | Additional seam equations; coupling with padding structure |
| 3 | Non-padded / long messages: does spectroscopy generalize beyond W[8..15] known? | Dependency window analysis for multi-block messages |
| 4 | Seam null space (36-dim, from Phase 1163) → GL(4,C) representation theory connection | Representation theory; prior NEXUS Phase 1163 results |
| 5 | GF(2) inversion of compression function rounds 1–6 ($Z_3$ hardness wall) | Combine schedule spectroscopy with $Z_3$ round inversion |
| 6 | BBP inverse: does address space cluster by pattern shape in the π-field? | Digit-address correlation study; parallel to schedule spectroscopy |

## 7. Summary

**The SHA-256 message schedule is structurally transparent.** The seam words W[16–30] are rotation-encoded address stubs: each word encodes the GF(2) image of a message word under σ0 or σ1, plus known or observable addends. Because σ0 and σ1 are rank-32 bijections over $GF(2)^{32}$, their inverses exist and can be applied exactly.

Given W[0] and the seam words, all 15 remaining message words are recovered exactly. Conditioned on seam transcript access, the 512-bit message-block reconstruction collapses to a 32-bit anchor search over W[0]. In ordinary preimage search those seam words are not observable from the digest; the SHA† open problem is whether H constrains compatible seam transcripts tightly enough to reduce the search. The compression function provides the second constraint layer that must be characterized.

Method: seam spectroscopy. Axis: orthogonal to execution (neither forward traversal nor backward reversal). Entry point: the seam K = {16,17} as established in prior NEXUS work. This is the side door — the axis the field was never braced against.

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com

Copyright Dean A. Kulik – Orcid ID # 0009-0003-3128-8828 – Ver Mark 9
Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0)
github.com/QuHarmonics/The-Nexus-Harmonic-Reality info@quharmonics.com