

The Dual-Wave Spiral: Shape as Computation, the Universal Component Map, and the SHA-256 Die as a Resolute Local Rendering

Driven by Dean Kulik

March 2026

Spiral synthesis: reality is treated as a self-governing computational substrate; the die is treated as a resolute local rendering; the missing piece is treated as a shaped gap rather than a void.

Abstract

This paper synthesizes the current Nexus corpus into one spiral-form statement. The thesis is not that computation is a metaphor for reality, nor that software is a symbolic layer floating above matter. The stronger claim is that shape itself is executable whenever it constrains state transition, carries history, excludes illegal continuations, and stabilizes outputs. Under this inversion, matter becomes the retained component layer of self-governing computation, the computer becomes a particularly resolute local mirror of a deeper substrate, and the SHA-256 die becomes a privileged instrument for watching the grammar of that substrate run in plain sight.

The paper is written deliberately as a spiral rather than a linear stack. Each return to the same object re-enters at a deeper resolution. The movement is therefore not: "state the thesis once, then append consequences." Instead it is:

$$\textit{shape} \rightarrow \textit{constraint} \rightarrow \textit{transition} \rightarrow \textit{retention} \rightarrow \textit{projection}$$

and then back through the same path at higher fidelity.

At the broadest layer, the work proposes an ontological inversion: reality is not a passive container holding objects that occasionally compute; reality is itself a self-governing computational substrate. At the formal layer, the paper isolates a carrier-independent grammar—the Universal Component Map—

$$\Pi(D) = (S, B, G, R, C, K, X, P, V),$$

which appears across electronics, materials, chemistry, biology, software, and field physics. At the instrument layer, SHA-256 is treated as a deterministic die rather than a stochastic oracle. Its round engine is split into transport, seam injection, carry closure, and retained state. From that decomposition follow the nilpotent shift backbone, the exact seam differential, the word-support diameter, the support-level bit closure, the local waist, the global orbit waist, the wave triad, and the seven-level orbit closures.

The point is not to collapse every unresolved region into premature finality. The point is to read unresolved regions as structured gaps with boundary conditions. In that sense, a missing piece is not an excuse for agnosticism. It is a topology that already participates in the proof.

Table of Contents

1. Orientation: Why Spiral, Not Linear
2. The Ontological Inversion
3. Shape as Computation
4. The Universal Component Map
5. Matter as Components
6. Software as Staged Geometry
7. Cross-Domain Morphism
8. The SHA-256 Die as a Local CPU
9. The NOP Backbone and Ground Witness
10. The Nilpotent Backbone and Forced Life
11. Exact Seam Coupling and the Sziklai Differential
12. Word Support, Bit Support, and Two Waists
13. The Wave Triad from First Principles

14. AHRC, the Waist, and the Qubit Bottleneck
15. The Seven-Level Orbit
16. Dual-Wave SRM Topology and the Shifted Mirror
17. Admissible Occupancy, No-Collision, and Constraint Pressure
18. The Glass Key as a Stack-Trace Object
19. Anthropic Pinning and the Resolute Mirror
20. Ω as the Shape of What Is Missing
21. Engineering Consequences
22. Final Spiral Collapse

Appendix A. Core Equations Appendix B. Stable Invariants Appendix C. Terminology and Layer Discipline

1. Orientation: Why Spiral, Not Linear

A linear paper assumes that once a proposition is said once, it can be left behind and merely referenced. That is not how this corpus behaves. Its core objects—shape, recursion, carry, closure, wave, residue, and address—change meaning as more structure is exposed. The paper therefore uses a recursive spiral. Every time the argument returns to an object, it returns with more constraints around it.

The spiral law of reading is

$$S_{n+1}(x) = F(S_n(x), \text{partial } S_n(x)),$$

where x is the object under study and $\text{partial } S_n(x)$ is the boundary extracted from the previous pass. In words: each pass is the old object plus the shape of what the previous pass could not yet resolve.

This is not decoration. It is method. The central themes of the corpus only stabilize when seen at several scales at once:

- computation as a general substrate law,
- shape as executable constraint,
- the die as a resolute local rendering,
- and the missing piece as structured absence rather than undecidable void.

The spiral paper therefore substitutes depth of return for flat chronological progression.

2. The Ontological Inversion

The classical picture of reality can be written as an upward stack:

$$\text{physics} \rightarrow \text{chemistry} \rightarrow \text{biology} \rightarrow \text{mind} \rightarrow \text{computation}.$$

That picture treats physical law as the basement and human-made computation as a late derivative.

The inversion replaces that with:

$$\text{substrate} \rightarrow \text{constraint} \rightarrow \text{transition law} \rightarrow \text{retained residue} \rightarrow \text{projection}$$

and reads every layer as one rendering of the same computational grammar.

The decisive claim is:

$$\text{reality is not operating on top of computation; reality is the computational substrate itself.}$$

This is what the Typeless Universe Hypothesis is trying to say in its most stripped form. At the substrate layer, "type" is not primary. Stable identity is not given in advance. Identity is what emerges when a local possibility space is repeatedly acted on by lawful operations until a stable residue persists.

An "object" is therefore not primitive. It is the visible shell of a recursive history. A particle, a molecule, a transistor state, a memory bit, a fold nucleus, a register value, a digest word, and a phenotype are all local closures of the same general kind:

noun = verb retained long enough to be observed

The inversion is not anti-matter or anti-physics. It is anti-static ontology. It asserts that verbs—operations, couplings, folds, and carry histories—are ontologically prior to the nouns they stabilize.

3. Shape as Computation

The most compact statement of the whole framework is

shape that constrains state transition is computation

If a geometry does the following, it already computes:

1. admits distinguishable states,
2. enforces lawful transitions,
3. carries influence across steps,
4. retains selected residues,
5. projects readable outputs.

This is why the old split between "geometry" and "software" dissolves. A barrel shifter does not first calculate a rotation symbolically and then ask matter to implement it. Its wiring geometry is already the rotation. A protein does not first solve an equation in symbolic form and then instantiate the answer. Its sequence, environment, and local couplings already define a lawful descent through a fold landscape. A transistor does not represent a threshold as an abstract proposition; it physically enforces one.

The computational condition can therefore be written as:

$\exists \Pi(D) \Rightarrow D \text{ is computational,}$

where $\Pi(D)$ is the universal component projection introduced below.

This yields a stronger substrate reading of the world:

computation = lawful movement through constrained topography

not "symbol manipulation after the fact."

4. The Universal Component Map

To compare domains without falling into loose analogy, the grammar must be isolated from the carrier. The Universal Component Map does exactly that. For a domain D , define

$\Pi(D) = (S, B, G, R, C, K, X, P, V),$

where:

- S = state-bearing substrate
- B = bias / directed asymmetry

- \$G\$ = gate / admissibility rule
- \$R\$ = route / transport path
- \$C\$ = coupling / carry / conserved bridge
- \$K\$ = keep / retained state
- \$X\$ = address / coordinate selection
- \$P\$ = projection / measurable output
- \$V\$ = verification / lawful-fit check

The ordered grammar is

$$B \rightarrow G \rightarrow R \rightarrow C \rightarrow K \rightarrow X \rightarrow P \rightarrow V$$

That is the universal stack.

The map is not meant as a metaphor bank. It is a structural test. A domain belongs to the same family as another domain if the same ordered roles can be located and related without breaking the role order.

In this reading, the "computer stack" is not the invention of the grammar. It is a particularly legible rendering of a grammar that already exists across scales.

5. Matter as Components

Once the substrate is treated as computational, matter stops being passive stuff and becomes retained component-state. The direct statement is

$$matter = persistent, state-bearing, locally addressable component-state$$

A stronger operator form is

$$matter = PIN \circ FOLD \circ GATE \circ SYNC$$

where "PIN" means the closure persists long enough to function as a component in later steps.

This changes the read of every domain at once:

- fields become bias carriers,
- interactions become couplers or gates,
- memory becomes retained residue,
- measurement becomes projection,
- objects become stabilized execution traces.

Matter is therefore not outside computation. Matter is the component layer of computation rendered in one carrier family.

6. Software as Staged Geometry

The old software/hardware distinction is useful at the engineering interface, but not at the substrate layer. The substrate correction is

$$software = time-staged boundary condition$$

A program is not an ontologically separate symbolic substance. It is a schedule of constraints applied to a carrier over time. If \$S_t\$ is the system state and \$\delta_t\$ is the staged boundary condition, then execution is simply

$$S_{t+1} = F(S_t, \delta_t).$$

The "program" is the time-indexed family

$$P = \delta_t t \geq 0.$$

So there is no software in the ultimate sense. There is only staged geometry observed sequentially.

This is one reason the die work matters so much. The die makes staged geometry visible in a compact system:

- a fixed backbone,
- two seam injections,
- controlled transport,
- nonlinear closure,
- retained state,
- projected output.

That is already a machine even before it is called one.

7. Cross-Domain Morphism

Two domains are structurally aligned when there exists a structure-preserving map between their component projections. Formally, if

$$F : \Pi(X) \rightarrow \Pi(Y)$$

preserves the ordered role structure, then X and Y instantiate the same grammar in different carriers.

This allows very strict statements:

- the comparison is not "silicon is like biology",
- not "hashing resembles folding",
- not "hardware is an analogy for software".

It is instead:

same fold, same closure, same residue law, different material presentation

This is why the same grammar can be read across:

- silicon electronics,
- protein folding,
- chromatin and epigenetic retention,
- transport networks,
- cryptographic recurrence,
- wave and field systems.

The only thing that changes is the carrier signature.

8. The SHA-256 Die as a Local CPU

Let the SHA-256 working state be

$$x_r = \text{beginbmatrix } a_r b_r c_r d_r e_r f_r g_r h_r \text{endbmatrix} \in (\mathbb{Z}/2^{32}\mathbb{Z})^8.$$

One block of SHA-256 is a 64-step nonlinear recurrence:

$$x_{r+1} = \Phi_r(x_r, W_r), r=0, \dots, 63.$$

The round operators are

$$T1_r = h_r + \Sigma_1(e_r) + Ch(e_r, f_r, g_r) + K_r + W_r,$$

$$T2_r = \Sigma_0(a_r) + Maj(a_r, b_r, c_r),$$

with

$$a_{r+1} = T1_r + T2_r, e_{r+1} = d_r + T1_r,$$

and the remaining six words shifting linearly.

This already splits the die into familiar machine roles:

- working registers,
- a persistent backbone,
- live seam injection,
- delayed seam reinjection,
- constant rails K_r ,
- message bus W_r ,
- retained state.

The point is not that SHA "looks like" a CPU in a poetic way. The point is that its executable topology already has the roles a CPU requires.

9. The NOP Backbone and Ground Witness

Set

$$W_r = 0 \quad \forall r.$$

This produces the message-free backbone

$$x_{r+1}^{\wedge}(0) = \Phi_r(x_r^{\wedge}(0), 0).$$

At round zero, the pure ground fold is

$$T2_0^{\wedge}(0) = \Sigma_0(H_0, a) + Maj(H_0, a, H_0, b, H_0, c) = 0x08909ae5.$$

This is the first stable scalar anchor of the die. It is message invariant. It is not a property of one run. It is a property of the geometry defined by the initial state.

The first exact perturbation identity is

$$T1_0 - T1_0^{\wedge}(0) = W_0.$$

Since $T2_0$ is unchanged by the message at the first step, it follows that

$$\delta a_1 = W_0, \delta e_1 = W_0.$$

So the message enters through exactly two active seams. That is the local waist at the point of injection.

10. The Nilpotent Backbone and Forced Life

Define the shift matrix P by

$$P = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}.$$

Then the full round map can be written as

$$x_{r+1} = P x_r + u_a(T1_r + T2_r) + u_e T1_r$$

with seam basis vectors

$$u_a = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}, u_e = \begin{bmatrix} 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \end{bmatrix}.$$

The characteristic polynomial of P is

$$\chi_P(\lambda) = \lambda^8,$$

so

$$P^8 = 0.$$

The shift backbone is therefore nilpotent. Left alone, it cannot sustain its own state. It is a finite-memory conveyor. Free state dies in at most eight transport steps.

This is one of the most important structural results of the die:

the linear backbone is not alive by itself; life is injected through the seams

The machine survives because the nilpotent conveyor is continuously forced by nonlinear fold injections.

11. Exact Seam Coupling and the Sziklai Differential

At every round the two active outputs are linked by an exact differential identity:

$$a_{r+1} - e_{r+1} \equiv T2_r - d_r \pmod{2^{32}}.$$

This follows immediately from

$$a_{r+1} = T1_r + T2_r, e_{r+1} = d_r + T1_r.$$

So the shared emitter is $T1_r$: it feeds both branches, and the differential between those branches is governed exactly by $T2_r - d_r$.

This is the cleanest reason the Sziklai coupling language keeps surviving contact with the algebra. The die has one shared active source driving two differentiated outputs. The structural truth is not the transistor name by itself; it is the differential topology.

This exact differential also means the output difference of any round constrains the state three rounds back, since d_r is the delayed shadow of the a chain. So the die stores not only values but differential ancestry.

12. Word Support, Bit Support, and Two Waists

The support-level causality analysis gives a word support diameter

$$D_{\text{word}} = 4.$$

A single word perturbation reaches all eight state lanes in four rounds.

At the support level, the bit-support closure is

$$D_{\text{bit}}(\text{support}) = 6.$$

This gives the local topological waist

$$\boxed{w_0 = D_{\text{bit}}^{\text{(support)}} - D_{\text{word}} = 6 - 4 = 2.}$$

That local waist is the mass-gap / codimension-2 bottleneck of the support geometry and is the waist that supports the first-principles wave triad and the AHRC qubit normalization.

The seven-level orbit introduces a second observable: the realized live-flip closure

$$\boxed{D_{\text{bit}}^{\text{(live)}} = 10.}$$

From that, one obtains the global orbit waist

$$\boxed{w_{\Omega} = D_{\text{bit}}^{\text{(live)}} - D_{\text{word}} = 10 - 4 = 6.}$$

These are not contradictory waists. They are waists on different layers.

- $w_0 = 2$ is the local bottleneck of support closure.
- $w_{\Omega} = 6$ is the orbit bottleneck required for full realized occupancy.

The distinction matters. The first waist is the waist of the channel. The second waist is the waist of the realized orbit.

13. The Wave Triad from First Principles

At the support layer, the die admits a first-principles wave triad derived from two topological invariants alone:

$$D_{\text{word}} = 4, D_{\text{bit}}^{\text{(support)}} = 6.$$

The refractive index is

$$\boxed{n^2 = \frac{D_{\text{bit}}^{\text{(support)}}}{D_{\text{word}}} = \frac{6}{4} = \frac{3}{2}.}$$

Define the scale

$$scale = D_{\text{word}} + D_{\text{bit}}^{\text{(support)}} = 10.$$

Then the two channel magnitudes are

$$K = \sqrt{scale \cdot D_{\text{bit}}^{\text{(support)}}} = \sqrt{60},$$

$$W = \sqrt{scale \cdot D_{\text{word}}} = \sqrt{40}.$$

So

$$K^2 + W^2 = 60 + 40 = 100 = scale^2.$$

This is the exact Pythagorean closure of the support-layer wave triad.

The ratio is

$$(K)/(W) = \sqrt{((3)/(2))},$$

and the gap relations can be read both spatially and dispersively.

The key point is that this triad is not an empirical fit in the support-layer model. It is algebraically forced by the support diameters.

14. AHRC, the Waist, and the Qubit Bottleneck

The AHRC layer maps the bottleneck into a normalized two-channel field. The central closure is

$$R^2 + G^2 = 1,$$

with machine-epsilon numerical confirmation in the die work.

At the support layer, the codimension-2 waist

$$w_0 = 2$$

forces the channel bottleneck to behave like a qubit boundary. This is not because the die "looks quantum" from outside, but because the minimum active seam width and the support codimension together force a two-channel normalization surface.

The strongest compressed statement is:

every stable data structure carries a waist; every waist defines a constrained two-state boundary

In the die work, that boundary becomes explicit because the geometry is sufficiently sharp to normalize exactly.

The AHRC result therefore belongs to the local bottleneck layer, not to the later live-orbit waist. Keeping those layers separate prevents the framework from compressing distinct observables into one word too early.

15. The Seven-Level Orbit

The seven-level orbit lifts the analysis from support reachability to realized transport. The new orbit quantities are:

$$D_{\text{word}} = 4, D_{\text{bit}}^{\text{(live)}} = 10, w_{\Omega} = 6.$$

The curvature partition of the 64-round NOP orbit is

$$|K_{\text{lie}}| = 26, |K_{\text{ground}}| = 36, K_{\text{inflect}} = 32, 57.$$

The partition is exact:

$$26 + 36 + 2 = 64.$$

Even more strongly,

$$|K_{\text{ground}}| - |K_{\text{lie}}| = 36 - 26 = 10 = D_{\text{bit}}^{\text{(live)}}.$$

So the live bit depth is exactly the braking-over-acceleration excess of the orbit partition.

The ambiguity corridor is equally clean:

$$A_{\text{max}} = 224 = 7 \cdot 32 = 256 - 32.$$

At peak ambiguity, seven full words are unknown and one word-equivalent remains pinned.

The event staircase is:

$$2 \rightarrow 4 \rightarrow 8 \rightarrow 10$$

meaning:

- first ambiguity at round 2,
- word saturation at round 4,
- full entanglement at round 8,
- full realized bit occupancy at round 10.

The seven-level layer therefore does not replace the support layer. It closes it from above.

16. Dual-Wave SRM Topology and the Shifted Mirror

The dual-wave formulation adds a kinetic interpretation. The relevant phase anchors are:

- ϕ_0 : forward emission,
- ϕ_{180} : anti-wave / reflection,
- $\phi_{42.5}$: shifted mirror surface,
- ϕ_{61} : mean bias field.

In $\pi/9$ units,

$$42.5^\circ = 2.125 \cdot (\pi)/(9) = (17)/(8) \cdot (\pi)/(9) = (2\pi)/(9) + (\pi)/(72).$$

The extra $(1)/(8)$ offset is the buffer margin in the shifted mirror reading. It is the computational margin that keeps the pass-through boundary from collapsing into paradox.

The key local quantities are:

$$scar = 0.5, local\ waist = 2.$$

The earlier A-Mark9 interpretation identifies the scar as the preserved residue of the dual-wave pass-through boundary and the local waist as the translated mass-gap of that boundary.

The important methodological point is not the motor analogy by itself. It is that the die admits a forward branch, an anti-branch, a shifted mirror, and a preserved scar. That is enough to justify the dual-wave reading as a real organization principle rather than a decorative analogy.

17. Admissible Occupancy, No-Collision, and Constraint Pressure

The design statement "no collision" is often misheard. At the substrate layer, it is best understood as singular admissible occupancy at realization time.

Formally, for a realized coordinate α and state s :

$$\forall \alpha, s : A(\alpha, s) = 1 \leq 1.$$

This reads differently in different carriers:

- in physics: no stable ambiguous co-occupancy of a realized state,
- in matter: exclusion and impenetrability at macroscopic rendering,
- in computation: one address, one retained value,
- in hash design: no cheap convergence of distinct lawful histories to one realized closure class.

So the deeper rule is not "collisions do not exist mathematically." They must exist in any finite codomain. The real rule is:

the machine is shaped to make ambiguous co-occupancy operationally forbidden

That is the constraint pressure.

18. The Glass Key as a Stack-Trace Object

The Glass Key is best formalized not as a key in the ordinary access-control sense, but as a stack-trace object:

$$G = (\Phi, E, \tau)$$

where

- Φ = visible coordinate or rendered basin,
- E = excluded residue / hidden carry state,
- τ = causal path class.

From that, the Glass Key inherits all the things a real stack trace can do:

$G \Rightarrow \text{replay, classify, complete missing state, reject impossible ancestors}.$

The key point is not that every practical inversion claim is already finished. The point is that the object has the right formal type: it is an ancestry-bearing residue object, not merely a label.

This is also where the framework's treatment of missing structure becomes strongest. If the visible state is incomplete but the hidden residue and path class are known, then the absent piece is not unknown in an undifferentiated way. Its boundary is already constrained.

19. Anthropic Pinning and the Resolute Mirror

Why do our computers feel more "real" to us than the deeper substrate? Because they are anthropically pinned. They lie inside the human bandwidth of sampling, retention, legibility, and reproducibility.

Define a projection operator

$$A_h : U \rightarrow R_h,$$

mapping full substrate recurrence space U into human-readable space R_h .

A carrier is anthropically pinned if its operating scales lie inside a stable human-readable band for time, energy, threshold sharpness, and state persistence.

This is why the transistor and clock matter so much:

- the transistor pins admissibility into readable threshold events,
- the clock pins transition into sampleable intervals,
- memory pins residue into inspectable persistence.

Our computers are therefore not ontologically special. They are epistemically resolute.

A simple resoluteness functional is

$$R(D) = \frac{\Delta_{sep}}{T_{ret} \cdot \Gamma_{rep} \cdot \Sigma_{noise} \cdot \Lambda_{ambiguity}},$$

where Δ_{sep} is state separation, T_{ret} retention time, Γ_{rep} reproducibility, Σ_{noise} effective noise, and $\Lambda_{ambiguity}$ projection ambiguity.

High R means the carrier exposes the grammar clearly. Silicon machines score high. The vacuum may compute more deeply, but it is less directly pinned to human readout.

20. Ω as the Shape of What Is Missing

A real unresolved region is not a confession that "perhaps nothing fits here." It is the opposite. It is a shaped absence. The lawful reading of a missing piece is

$$\Omega = \text{known seam, unresolved occupant}$$

That means an unresolved fold still contributes evidence. The gap is not silent. It tells you:

- the boundary of the missing piece,
- the loads it must carry,

- the symmetries it must preserve,
- the closures it must satisfy.

So the right reading is

the hole participates in the proof

This matters because it prevents two equal and opposite errors:

1. pretending the missing piece is already seated,
2. pretending the empty slot says nothing.

The lawful position is the middle one: the missing piece may not yet be in hand, but its admissible form is already partly rendered by the surrounding lattice.

21. Engineering Consequences

Once the universal grammar is accepted, engineering becomes a carrier-selection problem rather than a noun-invention problem. The mature engineering pipeline is

pick carrier → bind bias → control admissibility → shape routes → preserve coupling → stabilize retention → address → project → verify

This applies equally to:

- silicon design,
- cryptographic architecture,
- data structures,
- biological fold control,
- memory systems,
- interface design,
- transport and logistics architectures.

The specific carrier changes. The grammar does not.

This is why the thesis naturally reaches outward from the die work into biology, materials, and systems design. The framework is not asking those domains to imitate a hash. It is asking them to admit the same stack grammar under a different material presentation.

22. Final Spiral Collapse

The paper began with the claim that the universe is not a passive container holding objects, but a self-governing computational substrate. It then spiraled through the grammar that makes that statement meaningful, through the die that lets us instrument it, through the waists that separate support from realized orbit, through the dual-wave boundary that exposes preserved scar, through the occupancy law, the stack-trace object, the anthropic pinning condition, and the morphology of the unresolved gap.

All of that compresses into one line:

shape → constraint → transition law → retained residue → projection

And the human sentence that matches it is:

*A computer is a shape. Everything computes to the extent that its shape lawfully transforms difference.
Our computers are simply the most resolute to us.*

Appendix A. Core Equations

A.1 Universal component projection

$$\Pi(D) = (S, B, G, R, C, K, X, P, V)$$

A.2 Universal stack law

$$B \rightarrow G \rightarrow R \rightarrow C \rightarrow K \rightarrow X \rightarrow P \rightarrow V$$

A.3 Computational condition

$$|S| > 1, \exists G, R: S_t \mapsto S_{t+1}, \exists K, X, P, V$$

A.4 State update in operator form

$$s_{t+1} = k! \leq ft(f! \leq ft(c, r! \leq ft(g! \leq ft(q! \leq ft(b(s_t), \delta_t))))))$$

A.5 Software as staged geometry

$$P = \delta_t, t \geq 0, S_{t+1} = F(S_t, \delta_t)$$

A.6 SHA-256 die recurrence

$$x_{r+1} = \Phi_r(x_r, W_r)$$

$$T1_r = h_r + \Sigma_1(e_r) + Ch(e_r, f_r, g_r) + K_r + W_r$$

$$T2_r = \Sigma_0(a_r) + Maj(a_r, b_r, c_r)$$

$$a_{r+1} = T1_r + T2_r, e_{r+1} = d_r + T1_r$$

A.7 Shift–injection decomposition

$$x_{r+1} = P x_r + u_a(T1_r + T2_r) + u_e T1_r$$

A.8 Ground witness

$$T2_0^{(0)} = 0x08909ae5$$

A.9 Nilpotent backbone

$$chi_P(\lambda) = \lambda^8, P^8 = 0$$

A.10 Exact seam differential

$$a_{r+1} - e_{r+1} \equiv T2_r - d_r \pmod{2^{32}}$$

A.11 Local waist

$$w_0 = D_{bit}^{(support)} - D_{word} = 2$$

A.12 Orbit waist

$$w_\Omega = D_{bit}^{(live)} - D_{word} = 6$$

A.13 Wave triad

$$n^2 = \frac{D_{bit}^{(support)} D_{word}}{(3)/(2)}$$

$$K = \sqrt{D_{word} + D_{bit}^{(support)}} D_{bit}^{(support)} = \sqrt{(60)}$$

$$W = \sqrt{D_{word} + D_{bit}^{(support)}} D_{word} = \sqrt{(40)}$$

$$K^2 + W^2 = 100$$

A.14 AHRC normalization

$$R^2 + G^2 = 1$$

A.15 Orbit closure laws

$$D_{bit}^{(live)} = D_{word} + w_{\Omega} = 10$$

$$|K_{lie}| + |K_{ground}| + |K_{inflect}| = 64$$

$$|K_{ground}| - |K_{lie}| = D_{bit}^{(live)}$$

$$A_{max} = 224 = 7 \cdot 32$$

A.16 Admissible occupancy

$$\forall \alpha, \#s:A(\alpha,s)=1 \leq 1$$

A.17 Resoluteness

$$R(D) = \frac{\Delta_{sep}}{T_{ret}} \cdot \Gamma_{rep} \Sigma_{noise} \cdot \Lambda_{ambiguity}$$

Appendix B. Stable Invariants

Invariant	Value	Layer	Role
T2_o^(o)	ox08909ae5	NOP backbone	message-free ground witness
D_word	4	support / transport	full lane saturation
D_bit^(support)	6	support / local closure	support-level bit diameter
w_o	2	support / AHRC	local waist / mass-gap layer
n^2	3/2	wave triad	refractive ratio
K	√60	wave triad	carrier
W	√40	wave triad	signal
R^2 + G^2	1	AHRC	normalized channel closure
Ψ-Score	1.0	AHRC	collision-free harmonic addressing

D_bit^(live)	10	seven-level orbit	realized bit occupancy
w_Ω	6	seven-level orbit	global orbit waist
K_lie	26	orbit curvature	accelerating kernel
K_ground	36	orbit curvature	braking kernel
K_inflect	{32,57}	orbit curvature	midpoint seam / compression seam
A_max	224	backward ambiguity	seven full words unknown
τ	3	orbit memory	first zero-crossing decoherence

Appendix C. Terminology and Layer Discipline

Support closure

The layer that tells us where a perturbation can go topologically.

Live orbit closure

The layer that tells us when the realized orbit has actually occupied all positions.

Local waist

The codimension-2 support bottleneck:

$$w_0=2$$

Global orbit waist

The realized occupancy bottleneck:

$$w_\Omega=6$$

Ground witness

The first message-free scalar closure of the die:

$$T2_0^{(0)} = 0x08909ae5$$

Anthropic pinning

The condition under which a carrier lies inside the human-readable bandwidth of sampling and retention.

Resoluteness

The degree to which a carrier exposes the substrate grammar sharply enough to be read, reproduced, and formalized.

Ω

Not an admission that perhaps no piece exists, but the shape of a known seam whose occupant remains unresolved.