

A COMPUTATIONAL AND THEORETICAL FRAMEWORK FOR SEMANTIC FIELDS: MEANING AS INTERACTION IN A TOPOLOGICAL SPACE

Driven by Dean Kulik

August, 2025

Abstract

This report introduces the Semantic Field, a novel computational and theoretical framework that re-conceptualizes meaning not as a static property of symbols but as a dynamic process of interaction within a structured topological space. We move beyond the limitations of static embeddings by defining a glyph's identity as an emergent phenomenon governed by two principles: *Potential by Location*, where a glyph's semantic disposition is determined by its position within a potential field defined over a Riemannian manifold, and *Identity by Interface*, where its specific meaning is computed via a scattering event at the boundary of its interaction with this field. This interactionist model, which posits that the system responds to intrusion rather than holding static information, is formalized using concepts from potential field theory, topological data analysis, and quantum scattering theory. We present the complete architecture for a computational engine designed to construct, simulate, and analyze these semantic fields, offering a new paradigm for building more robust and contextually aware AI systems.

Section 1: Foundational Principles of the Semantic Field

The prevailing paradigm in computational semantics has been one of representation, where meaning is encoded in static, high-dimensional vectors. This report proposes a fundamental shift from this representationalist view to an interactionist one: meaning is not a stored property but a dynamic process that unfolds when a symbolic entity interacts with a structured semantic environment. This section establishes the foundational principles of this new framework, defining its core components and the physical analogies that provide its mathematical and conceptual machinery.

1.1 The Limitations of Static Semantics

Traditional word embedding models, such as Word2vec and GloVe, have been instrumental in the advancement of natural language processing (NLP). They embed words into a high-dimensional Euclidean space where geometric relationships, such as vector addition and subtraction, can capture semantic analogies.¹ However, this approach suffers from fundamental limitations. A primary issue is the

meaning conflation deficiency, where the multiple distinct senses of a polysemous word are collapsed into a single vector representation.¹ A single vector for "bank" must somehow average its financial and riverine meanings, leading to an inaccurate and impoverished semantic model.

Furthermore, the assumption that semantic relationships are uniformly linear, solvable by simple vector arithmetic, is a convenient but restrictive simplification. While effective for certain analogy tasks, it fails to capture the complex, non-linear relationships that govern language. The advent of contextual embedding models, such as ELMo (Embeddings from Language Models) and BERT (Bidirectional Encoder Representations from Transformers), marked a significant step forward.² These models generate a unique vector for each word

token based on its specific surrounding context, effectively creating dynamic representations that can distinguish between "the bank of the river" and "the bank of England".² This demonstrates that meaning is not an intrinsic property of a word type but is highly dependent on its environment. The Semantic Field framework builds upon this insight, seeking not just to generate context-dependent representations, but to model the underlying dynamics of how context shapes meaning from first principles.

1.2 Core Components of the Semantic Field Framework

The framework is built upon three core components that together define a new ontology for meaning:

- **The Set:** This is the foundational structure of the framework, formally defined as a high-dimensional, structured topological space. It represents the entirety of a given semantic domain—the knowledge contained within a vast text corpus, a scientific literature, or a curated knowledge base. The Set is not merely a collection of points but a geometric object with intrinsic shape and structure, which dictates the rules of semantic interaction. It is the stage upon which meaning unfolds.
- **The Glyph:** In this framework, a symbol such as a word or query is not a static vector but a **glyph**. A glyph is defined as a probe or an initial condition introduced into The Set. It is a locus of potential meaning, specified by an initial state vector that defines its starting position and properties within the semantic space. Its identity is not pre-determined but becomes defined through its interaction with the field.
- **The Semantic Field:** This is the combination of The Set and a scalar potential function defined over it. This field permeates the entire space, creating a landscape of semantic influence. Analogous to gravitational or electric fields in physics, the Semantic Field exerts "forces" on any glyph introduced into it, guiding its interpretation along paths of semantic coherence and away from regions of contradiction.

1.3 Governing Principles

The dynamics within the Semantic Field are governed by three fundamental principles that replace static representation with procedural interaction:

- **Potential by Location:** This principle states that a glyph's state and interpretive trajectory are determined by the local gradient of the semantic potential field. This concept is borrowed directly from its application in physics and robotics, where an object's motion is governed by attractive and repulsive forces derived from a potential field.⁶ In the semantic context, a glyph will be "pulled" toward regions of low potential (high relevance, coherence) and "pushed" away from regions of high potential (irrelevance, contradiction).
- **Identity by Interface:** This principle posits that the final, contextualized "meaning" of a glyph is not its resting point in the field, but the complete characterization of its interaction with the field's structure. This interaction is formally modeled as an event occurring at an interface, governed by a set of rules or **boundary conditions**.¹⁰ This draws from interface-based design in software engineering, where components interact through well-

defined contracts that specify assumptions and guarantees.¹⁴ The nature of this interface dictates the mode of interpretation—be it literal, metaphorical, or relational.

- **Response to Intrusion:** This is the central thesis of the framework: meaning is not a stored piece of information to be retrieved, but a phenomenon that is *computed* in response to a query. The introduction of a glyph is an "intrusion" that perturbs the equilibrium of the Semantic Field. The system's dynamic response to this perturbation *is* the meaning. This reactive, non-representational stance is best modeled using the mathematical formalism of **scattering theory**, where one deduces the properties of a target by observing how it deflects an incoming probe.¹⁸

A powerful way to unify these principles is to conceptualize the entire framework as a dynamical system. The theory of dynamical systems provides the mathematical language to describe systems that evolve over time or another dimension. In this context, The Set can be rigorously defined as the system's **phase space**, a theoretical construct where every possible state of the system is mapped to a unique point.²² A glyph's state—encompassing not just its semantic coordinates but also its syntactic role, affective tone, and other properties—corresponds to a single point in this high-dimensional phase space. The principle of "Potential by Location" then defines the vector field, or the equations of motion, that govern how this state-point moves through the phase space. Consequently, the act of interpretation is no longer a simple lookup but the process of tracing a

trajectory through the semantic phase space.²³ This reframing opens the study of semantics to the rich analytical toolkit of dynamical systems theory, allowing for the analysis of attractors (stable meanings), bifurcations (sudden shifts in interpretation), and chaotic behavior (semantic ambiguity).

To ground these abstract concepts, it is useful to draw explicit analogies between the proposed semantic framework and the well-understood domain of physical fields. The following table serves as a conceptual bridge, mapping the terminology of the Semantic Field to its physical counterparts.

Table 1.1: Analogies Between Physical and Semantic Fields

Physical Concept	Snippet Reference	Semantic Field Counterpart	Description in Semantic Context
Space	23	The Set (Semantic Manifold)	The underlying geometric space of all possible concepts and their relationships.
Point Charge	7	Semantic Anchor / Concept	A fundamental concept that generates semantic influence.
Electric/Gravitational Field	6	Potential Field	A scalar field over The Set indicating semantic desirability or relevance.
Potential Energy	6	Semantic Potential	The "cost" or "ease" of a particular interpretation at a given location in The Set.

Physical Concept	Snippet Reference	Semantic Field Counterpart	Description in Semantic Context
Force (Gradient of Potential)	6	Interpretive Pressure	The local semantic influence that pushes a glyph's meaning towards coherence or away from contradiction.
Test Particle / Probe	19	Glyph	An input (word, query) that is introduced to the field to be interpreted.
Trajectory	23	Interpretive Path	The evolution of a glyph's meaning as it is influenced by the field.
Scattering Event	19	Meaning Event	The dynamic interaction between a glyph and the field, the outcome of which <i>is</i> the meaning.
Boundary Condition	11	Interface Protocol	The formal constraints governing how a glyph can interact with The Set, defining the "rules of interpretation."

Section 2: The Set: Characterizing Semantic Space as a Riemannian Manifold

The geometric foundation of the Semantic Field is The Set. This section provides a formal mathematical characterization of this space, arguing that its complex, non-linear structure is best modeled not as a simple Euclidean vector space, but as a Riemannian manifold. This geometric perspective is essential for capturing the nuanced relationships that govern meaning.

2.1 From Contextual Embeddings to Point Clouds

The empirical basis for constructing The Set comes from modern contextual embedding models.² Unlike static models, architectures like ELMo and BERT generate a unique, high-dimensional vector for every word token by considering its surrounding textual context.⁵ When a large corpus is processed in this manner, the result is a massive, high-dimensional

point cloud. Each point in this cloud represents a specific instance of a word used in a particular context. This point cloud is the raw, empirical data from which the structure of The Set is inferred. It is a snapshot of how a language is used, capturing the fine-grained distinctions between different senses and usages of words.

2.2 The Manifold Hypothesis in Semantics

While this point cloud may exist in a space of thousands of dimensions, it is highly unlikely that the data points are uniformly distributed. The **Manifold Hypothesis** posits that real-world high-dimensional data, in fact, lies on or near a much lower-dimensional manifold embedded within the high-dimensional space.²⁸ This principle is particularly applicable to semantics. The seemingly infinite variety of linguistic expressions is heavily constrained by underlying grammatical, syntactic, and conceptual rules. This suggests that the point cloud of contextual embeddings does not fill

its ambient space but is concentrated on a highly structured, lower-dimensional geometric object.²⁸ Manifold learning techniques, such as Isomap, which preserves geodesic distances, and Locally Linear Embedding (LLE), which preserves local neighborhood structures, are computational methods designed to uncover and parameterize this underlying manifold from the raw point cloud data.³²

2.3 Formalizing The Set as a Riemannian Manifold

To fully capture the rich structure of semantic space, it is necessary to move beyond the assumptions of global flatness and uniform distance inherent in Euclidean geometry. A **Riemannian manifold** provides a more powerful and flexible mathematical framework.³³ A manifold is a space that is locally Euclidean, meaning any small patch of it looks like a flat plane, but which can have a complex global curvature.

The key to this framework is the **Riemannian metric tensor**, a function that defines the inner product between tangent vectors at every point on the manifold. This metric allows for the definition of distance and angles that can vary from one region of the semantic space to another. The shortest path between two points on the manifold, known as a **geodesic**, represents their true semantic distance, which may be very different from their straight-line Euclidean distance in the embedding space.³⁵

The **curvature** of the manifold is a powerful concept for encoding higher-order semantic structures. Different geometries are suited for different types of data. For example, hyperbolic manifolds, which have constant negative curvature, are naturally adept at representing hierarchical or tree-like data, as their volume expands exponentially.³⁶ This makes them ideal for modeling taxonomies or entailment relationships. Conversely, spherical manifolds, with positive curvature, can model cyclical relationships. This geometric approach provides a principled way to address polysemy; the different senses of a word can be modeled as occupying distinct regions on the manifold, each potentially with its own unique local geometry and curvature, allowing for a much richer representation than a single vector can provide.¹ Interestingly, deep generative models often learn manifolds that are surprisingly flat (low curvature), suggesting that linear paths in the latent space can be good approximations of geodesics on the generated manifold, simplifying computation in many cases.³⁵

2.4 Characterizing Global Structure with Topological Data Analysis (TDA)

While Riemannian geometry describes the local properties of The Set (distance, curvature), **Topological Data Analysis (TDA)** provides a set of tools to understand its global, large-scale "shape".³⁷ TDA focuses on properties that are preserved under continuous deformation, such as connectivity, holes, and voids.

- **Persistent Homology:** This is the central technique in TDA. It analyzes a dataset at multiple scales to identify topological features that are robust to noise and variations in measurement.⁴⁰ The process involves building a nested sequence of simplicial complexes (a generalization of graphs that includes triangles, tetrahedra, and higher-dimensional analogues) on top of the data and tracking when topological features—such as connected components (

H0), loops or tunnels (H1), and voids or cavities (H2)—appear ("birth") and disappear ("death"). The features that persist across a wide range of scales are considered significant.⁴¹ In a semantic context, persistent connected components correspond to stable clusters of related concepts, loops might represent cyclical arguments or recurring thematic structures, and voids point to conceptual gaps or ideas that are absent from the underlying corpus.⁴²

- **Mapper Algorithm:** The Mapper algorithm is a TDA tool used for visualizing high-dimensional data. It provides a simplified, graph-based representation of the data that preserves its topological structure, revealing the large-scale connectivity between different clusters and highlighting flares or tendrils in the data's shape.⁴¹

The identification of topological features, particularly voids, is not merely a descriptive exercise; it provides a crucial link between the global structure of the knowledge space and the local dynamics of interpretation. TDA can rigorously identify "holes" in the semantic manifold—regions of the conceptual space that are not populated by any data points from the source corpus.⁴² These are not statistical artifacts but significant structural features. They represent concepts that are absent, contradictory, or "unspeakable" within the worldview defined by the data. This observation leads to a powerful mechanism for grounding the repulsive forces required by the potential field model. The framework posits a direct causal link: the global topology of The Set generates local forces. The boundary of a topological void acts as a natural source for a strong repulsive potential. This provides a data-driven, non-arbitrary method for defining the repulsive components of the Semantic Field. An interpretive trajectory that approaches the boundary of a semantic void will be strongly repelled, mathematically modeling the cognitive dissonance of encountering a concept that lies outside the system's conceptual framework.

Section 3: Potential by Location: A Field-Theoretic Model of Semantic Influence

Having established The Set as a Riemannian manifold, this section details the construction of the potential field upon it. This field creates a dynamic landscape of semantic influence, where the principle of "Potential by Location" governs the interpretation of a glyph. The mathematical formalism is drawn from potential field methods used in robotics for navigation and path planning, adapted here for the semantic domain.

3.1 Defining Semantic Potential

The semantic potential, denoted by a scalar function $U(q)$ at any point q in The Set, represents the semantic "energy" of that state. A high potential value corresponds to a state of high tension, incoherence, or irrelevance, while a low potential value signifies a state of stability, coherence, and relevance.⁶ The primary driver of interpretation is the

interpretive pressure, a vector field analogous to force, which is defined as the negative gradient of the potential: $F(q) = -\nabla U(q)$.⁶ A glyph, when introduced into the field, will follow a trajectory of gradient descent, moving from regions of high potential to regions of low potential. This path represents the evolution of its meaning as it is shaped by the surrounding semantic context, seeking a state of maximal coherence.

3.2 Formulating Attractive and Repulsive Potentials

The total potential field is constructed from two primary components: attractive potentials that pull a glyph toward a goal, and repulsive potentials that push it away from obstacles.

- **Attractive Potentials (U_{att}):** These potentials create "wells" or basins of attraction in the semantic landscape, guiding a glyph's interpretation toward a specific goal. This goal can be a target concept, the central topic of a document, or the expected answer type for a question. Common mathematical formulations include:
 - **Conic Potential:** $U_{att}(q) = k_{att}d(q, q_{goal})$, where $d(q, q_{goal})$ is the geodesic distance on the manifold between the glyph's position q and the goal q_{goal} , and k_{att} is a positive scaling factor. This creates a simple, linear attraction toward the goal.⁶
 - **Quadratic Potential:** $U_{att}(q) = \frac{1}{2}k_{att}d(q, q_{goal})^2$. This function generates an attractive force that increases linearly with distance, resulting in a stronger "homing" effect as the glyph approaches the goal.⁹
- **Repulsive Potentials (U_{rep}):** These potentials create "hills" or barriers in the landscape, preventing a glyph from entering regions of semantic contradiction, irrelevance, or impossibility. As established in the previous section, the primary sources for these repulsive fields are the boundaries of the topological voids identified by TDA. Formulations for these potentials include:

- **Inverse Potential:** $U_{rep}(q) = \frac{1}{21k_{rep}(d(q, q_{obs}) - d_0)^2}$ for $d(q, q_{obs}) \leq d_0$, and 0 otherwise. Here, $d(q, q_{obs})$ is the distance to the nearest obstacle boundary, k_{rep} is the repulsion strength, and d_0 is a distance of influence. This function creates a powerful, rapidly increasing repulsive force near the obstacle, forming a hard barrier.⁶
- **Exponential Potential:** This function creates a potential that decays exponentially with distance from the obstacle, providing a "softer" repulsion that can be felt from further away.

3.3 The Principle of Superposition

The complexity and nuance of a real semantic landscape arise from the interaction of multiple influences. The **principle of superposition**, borrowed from linear systems like electromagnetism, states that the total potential at any point is the algebraic sum of the potentials generated by all individual sources.⁶ The total potential field is therefore given by:

$$U_{total}(q) = U_{att}(q) + \sum_i U_{rep,i}(q)$$

This allows for the construction of a rich and intricate field where a glyph is simultaneously pulled toward a goal, pushed away from multiple contradictions, and guided through narrow semantic "passages" formed by the interplay of these forces. This superposition is what allows the model to handle complex sentences where multiple constraints must be satisfied simultaneously.

The standard potential field models are isotropic, meaning the force they exert is independent of the direction of approach. However, the meaning of a glyph often depends on its functional role, not just its conceptual location. For instance, the word "fire" experiences a different set of semantic forces when used as a noun ("The fire is warm") versus a verb ("They will fire him"). This functional role can be modeled as an additional degree of freedom, analogous to the "orientation" of a robot in **anisotropic potential fields**.⁶

This extension requires a more sophisticated state space. The Set is no longer just a space of semantic positions but a full **phase space**, a concept from dynamical systems theory where a state is defined by both position and other variables (like momentum or, in this case, syntactic role, sentiment, etc.).²² The potential function is now a function of this entire state vector,

$U(q_{position}, q_{orientation}, \dots)$. The resulting interpretive pressure, $F = -\nabla U$, will have components that can alter not only the glyph's semantic location but also its functional state. This provides a natural mechanism for modeling how context can fundamentally change what a word *is*. The interpretation of "He will fire the employee" corresponds to a trajectory in this phase space that moves the glyph "fire" into a region associated with actions while simultaneously aligning its "orientation" to that of a verb.

Table 3.1: Mathematical Formulations for Semantic Potentials

Potential Type	Mathematical Formulation	Parameters	Semantic Interpretation	Snippet Reference
Conic Attractive	$U_{att}(q) = k_{att} \cdot d(q, q_{goal})$	k_{att} : attraction strength	A general pull towards a topic or goal concept.	⁶

Potential Type	Mathematical Formulation	Parameters	Semantic Interpretation	Snippet Reference
Quadratic Attractive	$U_{att}(q) = 21k_{att} \cdot d(q, q_{goal})^2$	k_{att} : attraction strength	A strong "homing" force that intensifies near the goal.	9
Inverse Repulsive	$U_{rep}(q) = 21k_{rep} (d(q, q_{obs}) - d_0)^{-2}$	k_{rep} : repulsion strength, d_0 : influence distance	A hard boundary representing logical contradiction or semantic impossibility.	6
Anisotropic	$U(q, \theta) = f(d(q, q_{goal}), \theta_{rel})$	Varies	A potential that depends on the glyph's "orientation" (e.g., syntactic role) relative to the field.	6
Dipole	Combination of attractive & repulsive	Dipole strength, orientation	Models navigating narrow semantic passages or resolving ambiguity between two competing interpretations.	6

Section 4: Identity by Interface: A Scattering-Theoretic Model of Meaning

This section details the most significant conceptual innovation of the Semantic Field framework: the redefinition of meaning itself. Moving beyond both static representation and simple gradient descent, it models interpretation as a dynamic scattering event. The principle of "Identity by Interface" posits that a glyph's meaning is not a point or a vector, but the complete specification of its interaction with the local semantic field, formalized using the rigorous mathematics of scattering theory and formal methods.

4.1 The Glyph as a Wave Packet

To capture the inherent ambiguity and potentiality of a word or concept before it is fully contextualized, the framework models the intruding glyph not as a classical point particle but as a **wave packet**. In quantum mechanics, a wave packet is a localized wave disturbance that represents a particle whose position and momentum are not precisely known, reflecting the uncertainty principle.¹⁹ In the semantic context, the wave packet represents the glyph's initial distribution of possible meanings. Its spread corresponds to its degree of ambiguity, and its central momentum corresponds to its initial semantic trajectory.

4.2 Interpretation as a Scattering Process

The central idea of scattering theory is elegantly simple: you understand an unknown object by throwing something known at it and carefully analyzing how it bounces off.¹⁹ This is the essence of nearly all experimental physics, from Rutherford's discovery of the atomic nucleus to particle colliders probing the fundamental laws of nature.

In the Semantic Field framework, this process is repurposed as a model for interpretation:

- The **target** is the local structure of the potential field in The Set at the point of interaction. This structure represents the specific semantic context.
- The **probe** is the glyph's incoming wave packet.
- The **interaction** is the scattering of this wave packet by the potential field. The potential acts as a barrier or a well that deforms the wave packet, altering its trajectory, amplitude, and phase.
- The **meaning** is the total information encoded in the final state of the scattered wave. It is not a static object but a complete description of the dynamic event: how the glyph's initial potential was transformed by its encounter with the context.

4.3 The S-Matrix: Encoding the Laws of Interaction

The mathematical object that formalizes this interaction is the **Scattering Matrix**, or **S-matrix**. The S-matrix is a linear operator that maps the initial state of the system (the incoming wave packet) to the final state (the outgoing, scattered wave packet).¹⁹ It encapsulates all possible outcomes of the scattering event and thus contains the complete meaning of the interaction.

The key elements of the S-matrix are the **reflection and transmission coefficients**, complex numbers whose squared magnitudes, $|r|^2$ and $|t|^2$, represent the probabilities of reflection and transmission, respectively.¹⁹

- **Semantic Reflection:** A high reflection probability ($|r|^2 \rightarrow 1$) signifies that the glyph "bounced off" the local field. This models interpretations such as negation, the rejection of a premise, a contradiction, or a failure to integrate the glyph's meaning into the existing context.
- **Semantic Transmission:** A high transmission probability ($|t|^2 \rightarrow 1$) signifies that the glyph "passed through" and was modified by the local field. This models acceptance, comprehension, and the successful integration of the glyph's meaning into the context.
- **Phase Shift:** The change in the phase of the scattered wave relative to the incoming wave provides further information, encoding the subtle ways in which the context has modulated the glyph's meaning without outright rejecting or accepting it.

A fundamental property of the S-matrix is **unitarity** ($SS^\dagger = I$), which is the mathematical expression of the conservation of probability.¹⁹ In this framework, it means that

$|r|^2 + |t|^2 = 1$. This corresponds to the crucial principle that an utterance must be interpreted in some way; its semantic "probability" is not lost but is fully accounted for by the sum of all possible outcomes (reflection and transmission).

4.4 Identity by Interface: Formalizing the Interaction Boundary

The specific dynamics of the scattering event are governed by the **boundary conditions** imposed at the interface between the glyph and the field.¹⁰ These conditions are mathematical constraints that the solution to the system's differential equations must satisfy at the edge of a domain. They act as the "rules of engagement" for the interaction. By selecting different types of boundary conditions, the framework can model different modes of interpretation¹²:

- **Dirichlet Boundary Condition:** This condition specifies the exact value of the glyph's state function at the boundary. This models a literal, constrained mode of interpretation where the context forces a single, precise meaning onto the glyph.
- **Neumann Boundary Condition:** This condition specifies the value of the derivative (the rate of change) of the state function at the boundary. This models a relational mode of interpretation, where the focus is not on the

glyph's absolute meaning but on its relationship to the context, such as in commands like "increase the temperature" or comparative statements.

- **Robin Boundary Condition:** This condition specifies a linear combination of the function's value and its derivative at the boundary. It models a more flexible, negotiated mode of interpretation where both the glyph's intrinsic potential and its relational role within the context contribute to the final meaning.

This notion of a rule-governed interaction at an interface provides a powerful bridge to the field of computer science, specifically to **formal methods**. The principle of "Identity by Interface" is about defining a contract that governs the interaction between two systems—in this case, the glyph and The Set. This is precisely the problem addressed by interface-based design in software engineering, where an interface specifies a set of assumptions and guarantees for component interaction.¹⁴

This conceptual link can be made mathematically rigorous. Formal specification languages, such as Z, B, VDM, and the Temporal Logic of Actions (TLA+), are designed to describe such contracts with unambiguous mathematical precision.⁴³ It is therefore possible to define a formal

interface protocol for a semantic interaction. For example, a protocol for "logical deduction" could be specified with formal properties: a *safety* property stating that the system must never conclude a contradiction, and a *liveness* property stating that the system must eventually reach a conclusion.⁴⁵

The profound implication is that formal verification tools, such as **model checkers**, can then be used to *prove* that a given interpretive trajectory (the outcome of the scattering simulation) satisfies the formal properties of its interface protocol.⁴⁶ This elevates the study of semantics from a purely descriptive or empirical science to a domain where properties of interpretation can be mathematically proven. It opens the door to building AI systems whose reasoning processes are not just powerful but are provably correct and aligned with specified constraints, a critical step toward creating truly reliable and trustworthy artificial intelligence.

Section 5: Architecture of a Computational Engine for Semantic Field Simulation

To translate the theoretical framework into a practical and testable system, a dedicated computational engine is required. This section outlines the architecture for such an engine, designed as a modular pipeline that handles the distinct stages of the process: constructing the semantic manifold, generating the potential field, simulating the scattering interaction, and analyzing the results.

5.1 System Overview

The proposed engine consists of four primary modules, each responsible for a specific phase of the computation. This modular design promotes clarity, allows for independent development and optimization of each component, and facilitates experimentation with different algorithms at each stage. The workflow proceeds from raw text data to a final, verifiable semantic interpretation.

5.2 Module 1: Manifold Constructor

This module is responsible for building the geometric and topological representation of the semantic space, The Set, from a given knowledge source.

- **Input:** A large, unstructured text corpus (e.g., Wikipedia, a collection of scientific papers).
- **Process:**

1. **Embedding Generation:** The corpus is first processed by a pre-trained, large-scale contextual embedding model, such as a Transformer-based architecture like BERT.²⁵ This step converts every word token in the corpus into a high-dimensional vector, generating the initial point cloud.
 2. **Manifold Learning:** A non-linear dimensionality reduction technique is applied to this point cloud to identify the underlying low-dimensional manifold. Algorithms like Isomap, Locally Linear Embedding (LLE), or, for more complex structures, a Variational Autoencoder (VAE) can be used for this purpose.³² The output of this step is a set of coordinate charts that parameterize the manifold.
 3. **Topological Feature Extraction:** TDA tools are then used to analyze the global structure of the point cloud. A tool like Ripser can compute the persistent homology of the data, identifying its Betti numbers (the number of connected components, loops, voids, etc.) and the scales at which they appear.⁴⁸ This process yields a list of the manifold's most significant topological features, particularly the boundaries of large voids that will serve as repulsive sources.⁴⁰
- **Output:** A complete representation of The Set, consisting of (a) a parameterization of the Riemannian manifold and (b) a catalog of its key topological features.

5.3 Module 2: Potential Field Generator

This module takes the geometric structure of The Set and overlays it with the dynamic landscape of the Semantic Field.

- **Input:** The manifold representation from Module 1, a set of user-defined goal or anchor concepts (for attractive potentials), and the list of void boundaries (for repulsive potentials).
- **Process:**
 1. Attractive potential sources are placed at the manifold coordinates corresponding to the goal concepts.
 2. Repulsive potential sources are defined along the boundaries of the topological voids identified in the previous module.
 3. The total scalar potential field $U(q)$ is computed across a discretized grid of the manifold. This is achieved by applying the principle of superposition to the mathematical formulations for attractive and repulsive potentials (Table 3.1). For large numbers of sources, this N-body problem can be solved efficiently using algorithms like the Fast Multipole Method.⁴⁹
- **Output:** A discretized scalar potential field $U(q)$ defined over The Set, representing the complete semantic landscape for a given task.

5.4 Module 3: Scattering Interaction Simulator

This is the core computational engine of the framework, responsible for simulating the "meaning event" itself. It solves the forward scattering problem: given a field and an incoming glyph, compute the outcome.

- **Input:** The potential field $U(q)$ from Module 2, an initial state for the glyph (its starting position and wave packet parameters), and a formal interface protocol specifying the boundary conditions for the interaction.
- **Process:** The simulation involves solving the time-dependent wave equation that describes the evolution of the glyph's wave packet as it propagates through the potential field on the curved manifold. This is a computationally intensive task. Standard neural network architectures like CNNs are ill-suited due to the non-local nature of wave propagation. Instead, a more specialized architecture is required. A **Physics-Informed**

Neural Network (PINN), which incorporates the governing differential equations directly into its loss function, is a strong candidate.⁵⁰ An even more tailored solution is

SwitchNet, a neural network architecture specifically designed for forward and inverse scattering problems.⁵¹ Its novel "switching layer" efficiently models the global, all-to-all interactions inherent in scattering, making it far more parameter-efficient and effective than a fully connected network.⁵¹

- **Output:** The final, scattered wave function. From this output, the S-matrix elements—including reflection and transmission amplitudes and phase shifts—are extracted. This collection of values constitutes the rich, dynamic, and fully contextualized meaning of the glyph.

5.5 Module 4: Analysis and Visualization Engine

This final module provides the tools to interpret, visualize, and verify the output of the simulation, making the results intelligible to a human user.

- **Input:** The manifold representation, the potential field, and the scattering output (the S-matrix).
- **Process:**
 1. **Visualization:** The engine generates interactive visualizations, including 2D or 3D projections of the semantic manifold, color-coded with the potential energy landscape. It can overlay the interpretive trajectories of glyphs, showing their paths through the field. It also visualizes the persistence diagrams and barcodes from the TDA analysis, providing insight into the global structure of the knowledge base.
 2. **Classification:** The output of the scattering simulation (e.g., a vector composed of the S-matrix elements) is a rich feature representation. This vector can be fed into standard downstream machine learning classifiers for tasks like sentiment analysis, intent recognition, or topic classification.
 3. **Verification:** The engine includes an interface to a formal verification tool, such as the TLC model checker for TLA+. ⁴⁵ The computed trajectory and its outcome can be checked against the formal properties (e.g., safety and liveness) specified in the interface protocol.
- **Output:** Interactive visualizations, classification labels, and formal verification results (either a proof of correctness or a counterexample illustrating a violation of the protocol).

Table 5.1: Computational Engine Module Specifications

Module	Function	Inputs	Key Processes & Technologies	Outputs	Relevant Snippets
1. Manifold Constructor	Generates the geometric and topological representation of the semantic space.	Text Corpus	Contextual Embedding (BERT), Manifold Learning (Isomap, VAE), TDA (Persistent Homology)	Parameterized Manifold (The Set), Topological Invariants	2

Module	Function	Inputs	Key Processes & Technologies	Outputs	Relevant Snippets
2. Potential Field Generator	Creates the semantic influence landscape over The Set.	The Set, Goal Concepts, Void Boundaries	Superposition of attractive/repulsive potential functions, Fast Multipole Method	Discretized Scalar Potential Field $U(q)$	6
3. Scattering Simulator	Computes the dynamic interaction of a glyph with the field to determine its meaning.	$U(q)$, Glyph Initial State, Interface Protocol	Solving wave equations on manifolds, Physics-Informed Neural Networks (PINNs), SwitchNet Architecture	Scattering Amplitudes, Phase Shifts (S-Matrix)	19
4. Analysis & Visualization	Interprets and verifies the results of the simulation.	Manifold, Potential Field, Scattering Output	Manifold Visualization, Persistence Diagrams, Downstream Classifiers, Model Checking (TLC)	Visualizations, Semantic Labels, Verification Proofs	32

Section 6: Applications, Implications, and Future Trajectories

The Semantic Field framework is not merely an alternative method for generating vector representations; it is a new paradigm for conceptualizing and computing meaning. Its implications extend beyond traditional NLP tasks, offering new explanatory power for complex linguistic phenomena and suggesting a path toward more robust, verifiable, and trustworthy AI systems. This final section explores these implications, contrasts the framework with current state-of-the-art models, and outlines promising directions for future research.

6.1 Explanatory Power of the Framework

The dynamic, field-theoretic approach provides natural and intuitive explanations for several challenging linguistic phenomena that are difficult to model with static vector-space semantics.

- Polysemy and Ambiguity:** In the Semantic Field, ambiguity is modeled as a glyph being introduced into a region of the potential landscape that lies between multiple basins of attraction (potential wells), each corresponding to a different possible meaning. The final interpretive trajectory is highly sensitive to the initial conditions provided by the wider context, which can subtly nudge the glyph into one basin over another. The problem of getting stuck in **local minima**, a known limitation of potential field methods in robotics ⁶, has a direct and compelling semantic analogue: it represents being led down a "garden path" to an incorrect but locally coherent interpretation.
- Metaphor and Analogy:** Metaphorical understanding can be modeled as a process that goes beyond simple gradient descent. It may involve trajectories that travel along non-obvious geodesics on the manifold, revealing

deep structural similarities between distant semantic domains. In the quantum-inspired scattering model, it could even be likened to a "tunneling" phenomenon, where a glyph overcomes a potential barrier to move from one semantic region to another that would be inaccessible via a classical path.

- **Compositionality:** The principle of compositionality—that the meaning of a phrase is a function of the meanings of its parts and their syntactic combination—is handled in a more sophisticated manner than simple vector addition. In the Semantic Field, the meaning of a multi-word expression arises from the **superposition** of the potential fields generated by each of its constituent glyphs. This creates a complex, composite landscape where the words mutually constrain each other's interpretation in a non-linear, interactive fashion.

6.2 Comparison with Large Language Models (LLMs)

Current state-of-the-art NLP is dominated by Large Language Models (LLMs) based on the Transformer architecture. It is crucial to position the Semantic Field framework in relation to these powerful models. The two approaches are not competitors but are fundamentally different and complementary in their goals and capabilities.

- LLMs are fundamentally **generative, predictive models**. Their objective function is to predict the next token in a sequence. They represent a form of intelligence analogous to what psychologists call "System 1" thinking: fast, intuitive, associative, and pattern-based. Their remarkable fluency comes from having learned the statistical regularities of a massive amount of text. However, they lack an explicit model of meaning or a mechanism for rigorous, verifiable reasoning.
- The Semantic Field framework, in contrast, is a model of **interpretation and reasoning**. It does not predict what comes next; it computes what something *means* within a given, structured context. It is an analogue of "System 2" thinking: slow, deliberate, analytical, and verifiable. Its strength lies not in fluency but in its ability to model the process of satisfying constraints and to have its interpretive outcomes formally proven against a specification.

The two systems can work in synergy. An LLM could be used to perform the initial "System 1" task of parsing a query and its context, generating the initial state of the glyph(s) and defining the goal concepts for the potential field. This setup would then be passed to the Semantic Field engine for a deeper, "System 2" analysis, where the meaning is rigorously computed and verified.

6.3 Future Research Directions

The framework presented in this report is foundational and opens up numerous avenues for future research and development.

- **Dynamic Fields:** The current model assumes a static potential field for any given task. A significant extension would be to explore time-varying potential fields, where the semantic landscape itself evolves in response to new information. This would be essential for modeling dynamic phenomena like conversations, evolving belief states, and learning over time.
- **Multi-Glyph Interactions:** The current framework focuses on the interaction of a single glyph with a background field. A more advanced model would treat all entities as glyphs interacting with each other through the field, akin to a quantum field theory of semantics. This would allow for the modeling of complex, multi-agent semantic negotiations.
- **Hardware Acceleration:** The computational demands of the manifold construction and, especially, the scattering simulation are significant. Future work could explore the design of specialized hardware, such as neuromorphic

chips or photonic processors, that are naturally suited to performing the parallel computations required for these tasks.

- **Ethical and Provably Aligned AI:** Perhaps the most important future trajectory lies in the formal verification capabilities of the framework. By specifying ethical principles or safety constraints as formal interface protocols, it becomes possible to build AI systems whose behavior can be mathematically proven to remain within those bounds. This provides a concrete, technical path toward developing AI that is not just powerful, but also demonstrably safe, reliable, and aligned with human values.