Walk Up from Behind: A Constraint Reads Its Own Residue

Driven by Dean A. Kulik

August, 2028

The geometry of a formula supplies a continuum and distinguishes nothing.
Everything that selects is arithmetic — and the same selection runs in silicon, in hashing, and in the digits of π.

Do not chase the named object. Ask what constraints make it unavoidable.
The route is the content; the object is the residue.

— working principle, NEXUS Framework

Nothing can change / everything must change. Held whole it has no consistent value. That is not a defect to resolve. That is the ground.

How to Read This Thesis

This document is arranged as a program, and it executes as you read it. Each part takes the output of the part before it as its input, applies one operation, and hands its residue forward. The operation is always the same operation. What changes is the substrate it runs on.

The instruction set is small. There are exactly four operations, and a fifth that is conspicuously absent:

```
FOLD      impose a fixed-point-free self-relation on a set
CLOSE     require that a thing which turns must return
PARTITION install a boundary; the residue is the observable
CARRY     retain what the partition drops (nothing is erased)


; there is no ERASE. That absence is the thesis.
```

The program does not begin by constructing anything. It begins with a single statement that has no consistent value, tries to satisfy it with one element, and fails. That failure is the first instruction, and every later part is the exception path it forces. By the end, the same four operations that failed on one element will have produced π, a working analog machine, a uniqueness theorem, a correction to how neural attention encodes position, and a claim about what the physical floor of computation actually charges for.

Three companion documents — the audit-trail papers on the void, the analog space, and the period relation — carry the full record of what was tried and killed along the way: seventeen retractions, each with the computation that ended it. This thesis does not repeat that record. It shows the working spine, states what is proved and what is not at every step, and draws the implications. Where a result is a theorem it is marked as one; where it is an interpretation or a conjecture it is marked as that and never promoted.

> *The reader who believes none of the framing can still run every computation in this thesis and obtain every number in it. That property — that the result survives disbelief in the story — is the point of writing it as a program rather than an argument.*

The Ledger of Claims

Every load-bearing statement in this thesis carries one of these grades. They are used strictly.

Grade	Meaning
THEOREM	Proved analytically. Stated with its assumptions and its proof.
REDUCTION	A rigorous restatement that moves a question to firmer ground.

Grade	Meaning
MEASURED	Established by computation at stated precision or coefficient height. Not a proof of non-existence.
INTERPRETATION	A reading of a theorem. May be false without falsifying the theorem.
CONJECTURE	Believed, with reasons. Open.
FALSIFIED	Tried, killed, and kept as a navigational marker.

Contents

Abstract

The Bailey–Borwein–Plouffe formula is taken here not as a result to be admired but as a specimen to be dissected — an instance of a general phenomenon in which a continuous structure supplies a family of possibilities and an arithmetic condition selects one member from it. The thesis establishes that phenomenon in one worked case with full rigor and then exhibits the same selection mechanism, computed, in three unrelated substrates.

The rigorous core is a chain of results about the integral representation of BBP. The numerator is shown to be a pole annihilator that cancels six of eight poles of its denominator, leaving one real pair and one complex pair; base sixteen and shell eight are shown to be not two facts but one, both read from a single pole at $\sqrt{2}\cdot e^{(i\pi/4)}$. The space of numerators whose integral is a rational multiple of π is shown to have dimension exactly two, and that two is shown to be rank–nullity on the reader's own basis rather than a property of π. Within that space a single arithmetic condition selects the BBP formula uniquely, and a Boundary Evaluation Identity moves that condition from eight coefficients onto two. The analog realization collapses from four precision-matched channels to four taps on one damped resonator, and the readout is proved to be a closed-form arctangent — which yields, as the central theorem, that among all single-pole harmonic read-heads on the π design curve, exactly one has an integer radix, and it is base sixteen with an eight-fold shell. The proof reduces by a Galois argument to five cases and needs no search.

The unifying claim is that the operation performing this selection — constraint-induced reduction of an admissible set — is substrate-independent. It is exhibited in the parity structure of even versus odd aggregation, where the same common-mode cancellation that makes the BBP logarithms vanish makes the uniform carrier vanish from an even XOR spread over GF(2); this is applied to a standing problem in neural attention, yielding a parameter-free correction to positional blindness. It is exhibited in the closure of the analog read-head against its own error, where the two elements needed to defeat the amplification are shown to be the machine's own discarded exhaust. And it is exhibited at the thermodynamic floor, where the Landauer cost is shown to be levied not per bit and not per value but per dropped coordinate, at the exact site a reader stopped tracking.

The implication drawn is ontological in scope but stated conservatively: if computation is the substrate-independent operation demonstrated here, then a physical system realizing a given operator exposes the invariant belonging to that operator, and the invariant belongs to the operator rather than to the substrate. The thesis is explicit about the boundary of this claim. It does not assert that any particular physical system implements the BBP read-head, and it separates, at every step, the theorems from the interpretations they invite.

Part I — The Ground: Two Prohibitions and Nothing Else

A program needs a first instruction. This thesis takes as its first instruction not a construction but a constraint that cannot be satisfied by the smallest thing one might offer it. Everything else is the exception path.

I.1 The base statement

Begin with a single statement, held whole:

nothing can change / everything must change

This is one statement, not two axioms. Held whole it has no consistent truth value, and that indeterminacy is not a defect to be resolved by splitting it into operands. It is the ground the program runs on. Splitting it into 'C0' and 'C1' and reasoning about the space between them turns a constraint into an idea; the discipline of this thesis is to keep it a constraint. Operationally it resolves into two prohibitions that a candidate structure must satisfy simultaneously:

```
HOLD   the structure must close on its own set
    (no element may map outside it — nothing escapes)
BREAK  no element may map to itself
    (nothing may hold still)
```

I.2 The first instruction fails

Offer the prohibitions a single element. Ask for a map on a one-element set that satisfies both. To break, the element must leave itself. To hold, it must land in the set. The set is only itself. No such map exists — not a trivial map, but zero maps. The exhaustive count over one element is empty.

This is the load-bearing failure of the entire thesis, so it is worth stating as a computation rather than a remark. Enumerate maps on an n-element set under both prohibitions:

```
n   BREAK only   HOLD only   BOTH
1       0           1         0    <- HALT: one cannot be
2       1           2         1    <- the only survivor
3       8           6         2
4      81          24         9

the BOTH column: derangements. at n=1 it is empty.
```

> *Existence does not begin with something appearing from nothing. It begins with nothing and with one both being inadmissible, and with two being the first arrangement the prohibitions permit. And at two there is exactly one of it: a pair trading places.*

I.3 What FOLD is

The survivor at n=2 is not two objects. It is one object with two positions — a 2-cycle, a single fold. The distinction matters and recurs throughout: the integer n counts positions, not things, and the things are the cycles. At n=4 the admissible arrangements are a fold-of-four or two folds; never four separate items. So the primitive operation is not 'add an element' — there is no second element to add, because there was never a legal first one alone. The primitive operation is FOLD: impose a fixed-point-free self-relation, and read the positions it forces.

A fold has a spectrum. A k-cycle's positions are the k-th roots of unity, and this is where geometry enters a program that so far contained none:

```
2-fold  roots {+1, -1}       one real axis: opposite and equal
3-fold  three cube roots      a complex pair, no real axis
4-fold  {1, i, -1, -i}       real axis AND imaginary: phase enters
8-fold  45° apart, radius set  the diagonal appears: √2 enters
```

Nothing was added to obtain the circle. No metric, no space, no time was assumed. The circle is what the prohibition permits once you ask for more than two positions, and the appearance of √2 at the eight-fold is the first quantitative fact the program produces on its own. Part III will show that this exact eight-fold, at this exact radius, is the pole structure of the BBP formula — and that base sixteen is the decay accumulated over one turn of it.

Part II — The Void: A Formula as the Residue of Prohibitions

The Bailey–Borwein–Plouffe formula is usually written as four channels with weights +4, −2, −1, −1 attached to the denominators 8k+1, 8k+4, 8k+5, 8k+6. This part shows that the four channels are a projection of a two-pole object, that the object is the eight-fold from Part I, and that the formula is not constructed but left behind — the residue after a small set of prohibitions has cut everything else away. The method is the one named in the epigraph: do not build the formula, build the void it is the only survivor of.

II.1 The integral representation and the denominator

Expanding the geometric factor 16^(−k) as an integral converts each shell sum into a definite integral over the unit interval. For a denominator family 8k+r,

$$S(r) = \Sigma[k\geq 0]\ 1/(16^k\,(8k+r)) = 16 \int_0^1 t^{(r-1)}/(16 - t^8)\, dt$$

and assembling the BBP combination with numerator exponent p = r−1 gives π as a single integral of a rational function with integer coefficients:

$$\pi = 16 \int_0^1 (4 - 2t^3 - t^4 - t^5)/(16 - t^8)\, dt$$

```
BBP series   = 3.141592653589793
BBP integral = 3.141592653589793  match: True
```

Every quantity in this integrand is algebraic — an integer-coefficient numerator over an integer-coefficient denominator, integrated between 0 and 1. There is no transcendental content anywhere in the parts list, and yet the output is transcendental. This is an engineering constraint before it is a philosophical one: the machine has no storage location for the constant it produces. π is not in the formula. It is what closing the formula costs.

Now read the denominator as an object rather than a divisor. $D(t) = 16 - t^8$ vanishes where $t^8 = 16$. All eight roots share the modulus √2 and differ only in phase:

```
tap  angle  pole                |pole|
 0     0   +1.414214 +0.000000j    1.414214
 1    45   +1.000000 +1.000000j    1.414214
 2    90   +0.000000 +1.414214j    1.414214
 3   135   -1.000000 +1.000000j    1.414214
 4   180   -1.414214 +0.000000j    1.414214
 5   225   -1.000000 -1.000000j    1.414214
 6   270   -0.000000 -1.414214j    1.414214
 7   315   +1.000000 -1.000000j    1.414214
```

> *This is the eight-fold of Part I, at radius √2. The equation $t^8 = 16$ is a closure condition: turn eight times, return, having decayed by sixteen. One turn, folded eight ways — that is the entire space the formula lives in.*

II.2 The numerator is a pole annihilator

Both polynomials factor, and — the observation the whole part turns on — they share factors:

$$16 - t^8 = -(t^2-2)(t^2+2)(t^2-2t+2)(t^2+2t+2)$$

$$4 - 2t^3 - t^4 - t^5 = -(t-1)(t^2+2)(t^2+2t+2)$$

Two of the three numerator factors are denominator factors. The rational function reduces from degree-five-over-degree-eight to degree-one-over-degree-four:

$$N(t)/D(t) = (t-1) / [(t^2-2)(t^2-2t+2)]$$

```
gcd(N, D) = t^4 + 2t^3 + 4t^2 + 4t + 4   -> annihilates 4 of 8 roots
reduced   = (t - 1)/(t^4 - 2t^3 + 4t - 4)
16 * integral = 3.141592653589793    match: True
```

Six of the eight poles cancel outright. The survivors are taps 0 and 4 — the real pair at ±√2, the 0°/180° axis — and taps 1 and 7, the complex pair at 1±i, the ±45° diagonal. The numerator's role, which looked like an arbitrary choice of weights, is exactly this: it selects which four of the eight fold-positions get read, and leaves the other four unwired.

> **Result II.1 (Pole annihilation). THEOREM.**
>
> The BBP numerator over $16 - t^8$ reduces to $(t-1)/[(t^2-2)(t^2-2t+2)]$, annihilating six of eight poles and leaving exactly one real pair and one complex pair. Verified as a polynomial identity.

Partial fractions on the reduced form give the object at its own rank — two arms:

Arm	Poles	Type	$\int_0^1$
A: $t / 4(t^2-2)$	±√2	real	−log(2)/8
B: $-(t-2) / 4(t^2-2t+2)$	1 ± i	complex	+log(2)/8 + π/16
Sum	—	—	π/16

The logarithms annihilate exactly — the real arm contributes −log(2)/8, the complex arm contributes +log(2)/8, and they cancel. The arctangent has no partner and survives alone. This cancellation is the first appearance of the parity mechanism that Part V will find again in a completely different substrate: an even pairing kills the common mode, and only the differential part carries through.

II.3 Base sixteen and shell eight are one fact

The surviving complex pole sits at $1+i = \sqrt{2}\cdot e^{i\pi/4}$. Expanding its factor as a power series gives a closed form for every coefficient, verified against the symbolic expansion through $j = 15$:

$$1/(t^2-2t+2) = \sum_{j\geq 0} t^j \cdot \sin((j+1)\pi/4) / 2^{(j+1)/2}$$

The phase factor is $\sin((j+1)\cdot 45°)$, which returns to its start after exactly eight steps. The amplitude factor decays by $\sqrt{2}$ per step, so over one full period the accumulated decay is $(\sqrt{2})^8 = 16$.

Constant	Read off	Value
shell	the pole argument	8
base	the pole modulus raised to the shell	$(\sqrt{2})^8 = 16$

Result II.2 (One location, two constants). THEOREM.

In the surviving complex arm, the shell length 8 is the period of the pole argument and the base 16 is the pole modulus raised to that period. Both are read from the single pole $\sqrt{2}\cdot e^{i\pi/4}$.

The phase factor vanishes at $j = 3$ and $j = 7$ — shell positions $r = 4$ and $r = 8$ — so the complex arm is structurally blind at the 180° and 360° marks, and the real arm covers precisely what the complex arm cannot reach. The division of labour between the two arms is not designed; it is forced by which directions each pole can address.

II.4 The void has dimension two — and the two is the reader's

Let the numerator range over all polynomials of degree at most seven — eight coefficients, one per shell position — and define the void as the set of numerators whose integral is a rational multiple of π. This is a property of the value, hence basis-independent. Exact rank computation on the six non-π transcendentals that can appear (log 2, $\log(1+\sqrt{2})$, log 3, log 5, $\arctan(1/\sqrt{2})$, arctan 2) gives rank six of eight, so the void has dimension two:

Vector	$n_0 \ldots n_7$	$16\cdot\int$	Σn_p
v_1	[−4, 0, 0, 2, 1, 1, 0, 0]	−π	0
v_2	[0, −8, −4, −4, 0, 0, 1, 0]	−2π	−15

v_1 is the BBP formula up to sign. Both verified by direct series summation to 30 digits and by 250-digit numerical integration. A subtlety here is worth the discipline it demands: the two is not a property of π. Read forward, the constraints cut eight dimensions to two. Read backward, the value line was always one-dimensional and the extra dimension is the reader's own redundancy. Testing across every denominator studied:

D(t)	void dim	pi formula?	kernel	1+kernel	match
16 - t^8	2	yes	1	2	True
64 - t^6	1	no	1	1	True
4 + t^4	1	yes	0	1	True
1 + t^2	1	yes	0	1	True
1 + t^8	1	yes	0	1	True

dim(void) = dim(π-image) + dim(zero-relation space)

Result II.3 (Rank–nullity on the reader). THEOREM (per construction), CONJECTURE (general form).

The void dimension decomposes as the dimension of the π-image plus the dimension of the zero-relation space. Base 16 / shell 8 has exactly one π-formula, the same as the others; what distinguishes it is a redundant tap — eight handles on a seven-dimensional value space. The extra dimension is the reader's, not π's.

This is the epigraph running on the formula itself. Never place the incompleteness in the object. The object had two poles and one line; everything else was addressing. The redundant dimension that looked like room in π was room in the reader's parameterization.

II.5 One arithmetic condition selects BBP

The void is two-dimensional, so membership does not select BBP. One further condition does, and it is the one visible in the coefficients: $+4 - 2 - 1 - 1 = 0$. For a general element $a \cdot v_1 + b \cdot v_2$ the coefficient sum is $-15b$, which vanishes exactly when $b = 0$. That single condition also excludes the zero-relation line $z = 2v_1 - v_2$, whose integral is exactly zero at 250 digits and whose coefficient sum is 15. DC-blocking does two jobs at once: it removes the informationless zero relation and cuts the remaining space to the unique BBP line.

Result II.4 (Boundary Evaluation Identity). THEOREM.

Let $D(t) = b + s \cdot t^n$ with $s = \pm 1$ and $b+s \neq 0$, and $\deg N < n$. Set $A = \gcd(N, D)$, $M = N/A$. Then $\Sigma\, n_p = N(1) = A(1) \cdot M(1)$, with $A(1) \mid (b+s)$ and $A(1) \neq 0$. Consequently $\Sigma\, n_p = 0 \Leftrightarrow M(1) = 0$.

Proof. $N = A \cdot M$ in $\mathbb{Q}[t]$; evaluate at $t = 1$. $\Sigma\, n_p = N(1)$ by definition. $A \mid D$ so $A(1) \mid D(1) = b+s$, nonzero by hypothesis, hence $A(1) \neq 0$ and the equivalence follows. The proof is one line; the content is where it puts the condition. DC-blocking is stated on N, which has eight coefficients; after annihilation it is a condition on M alone, and for BBP $M = 16 - 16t$ has two coefficients. The constant $A(1) = 15 = 3 \cdot 5$ is the boundary value of the two annihilated quadratics, and it divides $D(1) = 15$. The selection condition is written in coordinates the annihilation left behind.

Tested on seven void elements across five denominators; A(1) divides D(1) every time:

D(t)	n vector	value	Sum	A(1)	M(1)	ID
16 - t^8	[64,0,0,-32,-16,-16,0,0]	pi	0	15	0	True

```
16 - t^8  [0,64,32,32,0,0,-8,0]        pi   120   15   8  True
64 - t^6  [16,-24,-8,-6,1,0]         ZERO   -21   7   -3  True
 4 + t^4  [8,8,4,0]                 pi    20   5   4  True
 1 + t^2  [4,0]                   pi    4   1   4  True
 1 + t^4  [0,8,0,0]                 pi    8   1   8  True
 1 + t^8  [0,0,0,16,0,0,0,0]          pi   16   1   16  True
```

Scope, stated so it cannot drift: base 16 / shell 8 is the only construction in this set whose void has dimension two, hence the only one where DC-blocking leaves anything standing. 'Three prohibitions, one survivor' is what happens when the void has a spare dimension, not a universal architecture. That the spare dimension exists here, and only here, is not yet explained by the void construction — it is explained by the arithmetic theorem of Part IV.

Part III — The Machine: One Resonator, Read Four Ways

Part II left the formula as a two-pole object. This part builds the object as a physical instrument, and in doing so discovers that the instrument is far simpler than the formula suggested. The four-channel series demanded four precision-matched paths and multipliers to rebuild squared coordinates; the object it describes is a single damped rotation, and a damped rotation is one RLC tank. The reduction from four channels to four taps on one resonator is not an approximation — it is the recognition that the multipliers existed only to reconstruct a coordinate the resonator never leaves.

III.1 The readout is a closed-form arctangent

Take an underdamped RLC and sample it at $t_j = j\cdot\Delta t$. With one-step amplitude ratio $r = e^{-\zeta\omega_0\Delta t}$ and one-step phase advance $\theta = \omega_d\Delta t$, and choosing amplitude and initial phase so that $x(t_j) = r^{j+1}\cdot\sin((j+1)\theta)$, the harmonically weighted readout with $w_j = 1/(j+1)$ has a closed form. This is not numerical agreement; it is an identity from complex analysis, here reinterpreted as an analog readout law.

$$S = \sum_{m\geq1} r^m \sin(m\theta)/m = \mathrm{Im}[-\ln(1 - re^{i\theta})] = \arctan(r \sin\theta / (1 - r \cos\theta))$$

Result III.1 (Readout law). THEOREM.

The $1/(j+1)$-weighted sample sum of a damped rotation with pole $re^{i\theta}$ equals $\arctan(r \sin\theta / (1 - r \cos\theta))$. Verified against direct summation on six random (r,θ) pairs to 10^{-11}.

Output equals $\pi/4$ exactly when the argument equals one, which gives the design equation:

$$r (\sin\theta + \cos\theta) = 1$$

One equation, two unknowns — a curve of valid (r,θ) pairs, not a point. Every point on this curve produces π. This is the first appearance of a fact that becomes the thesis of Part IV: the geometry supplies a one-parameter family of π-producing machines and singles out no member of it.

III.2 The explicit component map

The design equation ties directly to physical component values through the standard series-RLC relations $\omega_0 = 1/\sqrt{LC}$ and $\zeta = (R/2)\sqrt{C/L}$:

```
r     = exp(-R dt / 2L)
theta = wo sqrt(1 - zeta^2) dt
OUTPUT = 4 arctan( r sin(theta) / (1 - r cos(theta)) )


pi setting:  zeta = 0.403712752   Q = 1.238504351
end-to-end:  OUTPUT = 3.141592653589794   exact: True
```

The resonator's z-plane poles land at $(1\pm i)/2$, which is exactly the reciprocal conjugate of the BBP t-plane poles $1\pm i$ — the inside/outside duality of a geometric series. The RLC is not an analogue of the BBP object. It is the BBP object's surviving pole, instantiated.

Result III.2 (Pole correspondence). THEOREM / INTERPRETATION.

The resonator's z-plane pole is 1/conj(t-plane pole): $(1+i)/2 = 1/\mathrm{conj}(1+i)$, verified to 60 digits. That the two are therefore 'the same object' is an interpretation the identity invites but does not compel.

III.3 Four taps, and the balance that is not trimmed

The surviving structure in reduced form carries three logarithmic conditions and one arctangent:

$$16\int_0^1 = 4\pi(c+d) - 8(a+c)\cdot\log 2 - (16b/\sqrt{2})\cdot\log(1+\sqrt{2})$$

Condition	Algebra	Realisation	Trimmed?
kill $\log(1+\sqrt{2})$	$b = 0$	no DC path into real arm	no — an absence
kill log 2	$a + c = 0$	taps 0/4, 180° apart	no — differential pair
DC-block	$d = 2a$	2:1 ratio, centre tap	no — power of two

Taps 0 and 4 are exactly antipodal, so $a + c = 0$ is what one obtains by taking both ends of a single differential line. The condition that was derived algebraically in Part II as 'antisymmetric wiring' is not a technique to be tuned; it is the geometry of the ring. The balance is enforced by the four taps not wired, not by matched components.

```
phase per tap     45.0 deg
amplitude per tap  0.707107  = 1/sqrt(2)
Q           1.238504
over 8 taps: (1/sqrt2)^8 = 0.0625 = 1/16


rebuild pi from tap series: sum = 0.785398163397 = pi/4  (match: True)
```

> *The precision wall of the original four-channel design — four matched paths, one-percent resistors, a cancellation ratio of six — was an artifact of building the series rather than the object the series describes. The object is one heavily damped tank, and its balance is topological.*

III.4 The machine defeats its own error with its own exhaust

An honest open-loop read-head has a real defect: each ×16 recurrence stage that advances the digit position also amplifies error, with Lyapunov exponent ln 16. A perturbation ε becomes 16ε in one step and $16^k\cdot\varepsilon$ after k, so the open machine can read only low positions. The closure is the part of this thesis where the framework's method — look for what has been discarded — pays a concrete engineering dividend.

Two elements are needed to defeat the amplification, and prior analysis of scale-invariant readouts (the companion SILR work) fixes their requirements exactly: a lattice reference for what the true answer is, and a fixed noise ruler that does not adapt to the growing error — because a scale-invariant correction normalizes the ×16 growth away and goes blind to it. Both are already present in the machine, currently discarded to ground.

Needed	Provided by	Mechanism
a lattice reference	the reset comb	quotient counts, quantized by construction
a fixed noise ruler	the bounded ramp height	the fixed height IS the ruler
the snap	reset comparator	injection locking, not a digital step
the timing	settle-lag to the lattice	self-clocked, no external clock

Result III.3 (Closure by exhaust). MEASURED.

> The reset comb supplies the integer lattice and the bounded ramp supplies the fixed noise ruler; the snap onto the lattice is injection locking performed by the reset comparator already present. Closed, the loop reads arbitrary positions up to the component noise floor. Every claim verified by live computation in the companion paper.

The organizing principle stated there is that the gap between fixed points is a gradient the machine runs on: it drives the snap, it sets the timing, and it carries the carry. And the noise floor itself, the value 0.5, is not a tolerance to be out-run — it is the rounding operator itself, the boundary between two lattice cells. This is the same object as Part I's separatrix and Part V's partition, appearing now as the physical floor of an analog computation: the gap is where the digital lives.

Part IV — The Crown: Geometry Offers a Continuum, Arithmetic Selects One

Part III established that every point on the design curve r(sinθ+cosθ)=1 produces π. This part asks which of those points is base sixteen with an eight-fold shell — and proves that exactly one is, that the selection is entirely arithmetic, and that the geometry which produces π everywhere distinguishes nothing. This is the central theorem of the thesis, and its proof is the clean form the framework was reaching for: walk up from behind, ask what makes the object unavoidable, and find that the route — not the object — is the content.

IV.1 π belongs to no point of the curve

First, the negative fact that reframes everything. Fix the reader as w = 1/m and move the pole along the design curve. Every point returns π:

```
theta(deg)   r      4*Im(-log(1-z))
 22.5   0.765366865   3.141592653589793
 30     0.732050808   3.141592653589793
 36     0.716719836   3.141592653589793
 45     0.707106781   3.141592653589793
 60     0.732050808   3.141592653589793
 72     0.765366865   3.141592653589793
```

And every other reader reaches π too, at a different pole: the dilogarithm weight $1/m^2$ reaches it at r = 0.831983, the trilogarithm at r = 0.935992, the exponential at r = 0.703916. π is a level set that every (dynamics, reader) family crosses. It is not privileged by the operator at all.

> **Result IV.1 (π is a level set). MEASURED / THEOREM.**
>
> The observable belonging to a linear pole and a weight sequence is Im G(z), the reader's generating function at the object's pole. π is the value this takes along an entire curve of poles for any fixed reader, and along a curve of readers for any fixed pole. It belongs to neither the dynamics nor the reader in isolation, and it distinguishes no member of either family.

IV.2 The arithmetic condition

What singles out base sixteen is not π but digit extraction, and that imposes two discrete conditions the geometry knows nothing about. For the read-head to be a base-b digit machine the shell must close, θ = 2π/n for integer n, and the radix must be an integer, r^n = 1/b. From the design curve $r^2 = 1/(1 + \sin 2\theta)$, so

$$b(n) = (1 + \sin(4\pi/n))^{(n/2)} \quad \text{must be an integer} > 1$$

Convergence ($r < 1$) with shell closure forces $\sin 2\theta > 0$, hence $n \geq 5$. The theorem is that $b(n)$ is an integer greater than one for exactly one admissible n.

The proof, in the order it was found and then repaired

The first attempt asserted that integrality of b forces $\sin(4\pi/n)$ rational and invoked Niven's theorem. That inference is false: $\sqrt{2}$ is irrational and $(\sqrt{2})^4 = 4$ is an integer, so a power can be integral without the base being rational. The repair is a Galois argument, and it is the load-bearing step.

> **Lemma IV.2 (Integrality forces rationality). THEOREM.**
>
> Let $s = \sin(4\pi/n)$. Every Galois conjugate of s is a value $\sin(4\pi k/n)$, hence real and in $[-1,1]$, so $1+s^\sigma \in [0,2]$ for every σ. If $b \in \mathbb{Z}$ then b is Galois-fixed, so $(1+s^\sigma)^{(n/2)} = (1+s)^{(n/2)}$ for every σ (apply to b^2 for odd n). Since $t \mapsto t^{(n/2)}$ is strictly increasing on $[0,\infty)$, two non-negative reals with equal positive powers are equal, so $1+s^\sigma = 1+s$ for all σ, hence $s \in \mathbb{Q}$.

Only now does Niven's theorem apply: a rational sine at a rational multiple of π lies in {0, ±1/2, ±1}. Five cases remain, and enumerating them requires no search:

s	n	b = (1+s)^(n/2)	Verdict
1	8	16	ADMISSIBLE
1	8/5, 8/9, 8/13	—	n not an integer
1/2	24	531441/4096	b not an integer
0	4	1	$n < 5$, and $b = 1$
−1/2, −1	—	—	$\sin 2\theta < 0$, violates convergence

> **Result IV.3 (Uniqueness of the integer radix). THEOREM.**
>
> Under the assumptions above — harmonic weights on a single damped pole pair, observable π, closed shell, integer radix, convergent — the only admissible configuration is $\theta = \pi/4$, $n = 8$, $r = 1/\sqrt{2}$, $b = 16$. The failure elsewhere is not marginal: at $n = 12$, $s = \sqrt{3}/2$ and its conjugate $-\sqrt{3}/2$ give $(1+s)^6 = 42.218744$ against $(1-s)^6 = 5.78e{-6}$, so $b(12)$ is not even rational.

IV.3 What fell away, and why that is the point

Earlier forms of this argument leaned on machinery that the final proof does not need: a monotonicity lemma for $b(n)$ with the clean derivative $h'(x) = x/(1+\sin x)$; an explicit bound $b(n) < e^{(2\pi)}$; the transcendence of $e^{(2\pi)}$ by Gelfond–Schneider; and a finite exhaustion over 42,970 shell counts. All four are correct. None is required. The Galois step removes the search entirely, and the tendency of the proof to shed heavy external machinery as it matured is the usual sign that the argument has become intrinsic to its object rather than borrowed.

> *The geometry supplies a smooth one-parameter family and distinguishes no point on it. Every constraint that selects is arithmetic — shell closure, integer radix, convergence — and the one point that is simultaneously on the π-locus and a digit machine is picked out by an integrality condition the resonator knows nothing about. That is the whole thesis, proved in one case.*

This also explains the loose end left in Part II.5: base sixteen was the only construction whose void had a spare dimension. That is the same fact as this theorem, seen from the coefficient side. The redundant tap, the dimension-two void, and the unique integer radix are three views of one arithmetic coincidence — that $\sqrt{2}$ lands the eight-fold shell inside the field $\mathbb{Q}(\zeta_8)$, where the eighth roots of unity and $\sqrt{2}$ generate the same structure.

Part V — The Same Operation Elsewhere: Parity, Aggregation, and Attention

Everything so far has concerned one formula. If the selection operation — an even aggregation annihilating the common mode, leaving only the differential structure — is genuinely substrate-independent, it must appear where BBP does not. This part exhibits it in the linear algebra of the SHA-256 diffusion functions over GF(2), and applies it to a standing problem in neural attention. The section is written with particular care about what is and is not established, because a memory of this work carried an overstated version of the central claim, and the overstatement is instructive.

V.1 The claim that had to be killed

A prior summary asserted a clean parity law for the SHA-256 Σ functions: that every 2-copy XOR spread of them has rank 31 and every 3-copy has rank 32. Direct enumeration — run before restating — falsifies it. Among pairs of the four sigma functions, the pair (σ_0, σ_0-lower) reaches full rank 32, and among triples, (Σ_1, σ_0, σ_1) drops to 30. The specific sigma functions do not obey a copy-count parity law.

```
2-copy sigma pairs:  ranks 31, 31, 31, 31, 31, 32   <- (0,2) breaks it
3-copy sigma triples: ranks 31, 31, 32, 30        <- not all 32
'every 2-copy rank 31, every 3-copy 32' : FALSIFIED
```

Result V.1 (Sigma parity, corrected). FALSIFIED, then narrowed.

The blanket copy-count rank law for the SHA-256 sigma functions is false by enumeration. What survives is a precise statement about the carrier, established next. Recording the falsification is not incidental: it is the difference between a claim that can be refuted in thirty seconds and one that holds.

V.2 What is actually true — the carrier is the all-ones vector

The defensible result concerns the uniform direction specifically. The map $\Sigma_0 \oplus \Sigma_1$ over $GF(2)^{32}$ has a one-dimensional null space, and its sole null vector is exactly the all-ones word:

```
null space of (Sig0 XOR Sig1): dimension 1
sole null vector: 0xFFFFFFFF
is all-ones: True
```

And the carrier parity that the specific sigma functions failed to obey does hold cleanly for the operation the aggregation actually performs: k identical copies of a value at fixed rotational offsets, XORed. The uniform direction appears in every copy, so its image is (k mod 2) times itself over GF(2) — it survives iff k is odd. Verified with real bit-shifts:

```
offsets [0,7]       k=2  rank 31/32   carrier ANNIHILATED
offsets [0,7,18]     k=3  rank 32/32   carrier SURVIVES
```

```
offsets [0,7,18,3]    k=4  rank 30/32   carrier ANNIHILATED
offsets [0,7,18,3,11]  k=5  rank 32/32   carrier SURVIVES
```

Result V.2 (Carrier parity). MEASURED.

$\Sigma_0 \oplus \Sigma_1$ has a one-dimensional null space spanned by 0xFFFFFFFF. For k identical rotated copies aggregated by XOR, the uniform carrier survives iff k is odd. The rotation-based spread obeys the parity law exactly; the specific sigma-function version does not.

> *This is the BBP cancellation in a different substrate. In BBP, the two real poles form an even pair and their logarithms — the common mode — cancel, leaving the differential arctangent. Over GF(2), an even aggregation sends the uniform carrier into the null space, leaving the differential structure. Same parity mechanism: even kills the common mode, odd preserves it.*

V.3 Application — attention's positional blindness is a null-space fact

Scaled dot-product attention is permutation-equivariant: reordering input tokens leaves the output unchanged. The field repairs this by injecting positional information from outside — sinusoidal, learned, or rotary encodings — all treating position as metadata attached to the signal rather than structure recoverable from it. The parity result suggests a different decomposition, in which the routing stays learned but the collision becomes a fixed, parameter-free spread.

The mechanism of the blindness is elementary and, stated as a design constraint, apparently unremarked: a weighted sum is a linear map from n dimensions to one, so by rank–nullity it has an (n−1)-dimensional null space. Everything in that null space is invisible to the aggregation. Uniform attention weights place the permutation-distinguishing information precisely there.

Configuration	Aliasing	Learned params
standard attention, uniform weights	100% (all 120 orderings → one output)	—
with hash-geometric spread channel	3.4%	zero

Result V.3 (Spread corrects aliasing). MEASURED.

On a 5-token permutation test, standard attention under uniform weights aliases all 120 orderings to a single output; a parameter-free hash-geometric spread channel reduces aliasing to 3.4%. The rank of the combined value-plus-spread map grows linearly, adding d_v dimensions per independent offset to full rank.

The falsifiable prediction, stated as such: on computed (non-random) data, the spread channel recovers positional information that standard attention cannot access without external encoding. The claim is structural, not a benchmark improvement — that experiment has not been run, and the thesis does not assert it. What is claimed is that positional blindness is a parity property of even-arity aggregation, that the spread corrects it at the level of the null space, and that the correction requires no learned parameters. Position was never metadata to be injected; it was structure the aggregation had projected away, recoverable at another address.

V.4 The pattern, stated at the right altitude

Three substrates have now shown one operation: the pole annihilation of a rational function, the null space of a GF(2) map, and the surviving attractors of a constrained dynamical system (the phase-oscillator computation of Part I's companion analysis). These are different objects and the thesis does not claim they are the same object. It claims they instantiate one abstract operation — constraint-induced reduction of an admissible set — and that this operation, not any of its instances, is the invariant worth naming.

The honest form of the cross-domain claim is a schema with stated content, not a universal identity. For a linear evolution with pole z read by a weight sequence, the observable is Im G(z), the reader's generating function at the object's pole. This is non-vacuous exactly on linear time-invariant systems with weighted readout — it holds for BBP, resonators, PLLs, and tapped filters, and it demonstrably excludes SHA-256, protein folding, and nonlinear oscillator networks, which do not admit a generating-function readout. That exclusion is the useful output: it converts 'the same architecture is everywhere' from an assertion into a specific debt, namely what would have to be shown to extend the schema to a nonlinear evolution.

Part VI — The Stack, the Toll, and What the Claim Actually Is

The thesis has run its program: a failed first instruction, a fold, a closure yielding π, a machine, a uniqueness theorem, and the same operation in three substrates. This final part draws the implications, builds the layer structure they imply, and states the ontological claim at the altitude the evidence supports — which is lower than the framework's ambition and is stated that way on purpose.

VI.1 The stack: carrier, partition, residue

Every result in the thesis has the same shape, and the shape composes. A carrier that cannot hold still (Part I's prohibition) is partitioned (a boundary is installed), and the residue of that partition becomes the carrier for the next layer. Analog is the carrier alone — any wave, no schema, no bits; the typeless ground. Digital begins the instant the first partition is installed, and the partition is never in the carrier — it is held by the reader.

Layer	Carrier	Partition	Residue
0 → 1	motion, free orbit	the gaps (separatrices)	bits (basin labels)
1 → 2	bits	couplings (constraints)	Fix(D|C), survivors
2 → 3	answers	boundaries	nesting, amplification

Three invariants hold at every layer, and each was measured in the course of the thesis. The partition is never in the carrier: it travels separately, as the reader. Each layer satisfies the change-mandate only by resting on the one below it — no layer is self-supporting. And going up amplifies while breaking the sum: nesting is a nonlinear gain, which is why a two-percent edge per point becomes a twenty-percent edge per match, and why the levels do not commute with addition.

> *Going up the stack is forgetting, not building. Each layer knows strictly less than the one beneath it, which is why the codebook always travels separately — it is the section of a quotient, and a quotient has no inverse without one. The reader is not observing the message. The reader is holding the half that was quotiented away.*

VI.2 The interface is a section; the leak is the fibre

The relationship between adjacent layers is a quotient and its section. A quotient π forgets; a section s chooses a representative. The asymmetry $\pi(s(y)) = y$ always, while $s(\pi(x)) = x$ almost never, is the entire structure of an abstraction. A quotient with a non-trivial fibre is precisely an abstraction that hides something, and the fibre is measurable: three implementations of the identical interface 'return the sum' agreed to nine decimal places while their timings spread by 151×. Time is in the fibre.

Result VI.1 (Leak is structural). REDUCTION.

> A zero-leak abstraction is the identity map: fibre one everywhere means bijection means nothing forgotten. Hence every non-trivial abstraction leaks, and a side-channel is a section an adversary constructs to read a fibre coordinate the designer believed was quotiented away. The leak is not a defect; it is the evidence that something runs underneath, which is the change-mandate seen from above.

VI.3 The toll: what Landauer actually charges for

At the thermodynamic floor the thesis makes its sharpest correction to received framing. The Landauer cost kT·ln2 is not charged per bit, not per value, and not per logical operation. Erasing a bit to 0 and erasing it to 1 cost the same; writing a known value over a known value costs zero; NOT, SWAP, and CNOT cost zero because they are bijections that compress nothing. What is charged is local phase-space compression — and the bill is levied on the reader's uncertainty about a specific site.

```
gate    cost / kT ln2     what it compresses
NOT        0        nothing (bijection)
SWAP       0        nothing (bijection)
CNOT       0        nothing (bijection)
XOR     1.000       two inputs -> one
AND     1.189       two inputs -> one (biased)


same site, erase cost depends on the READER:
knows the bit exactly        0.0000
knows it is 0 with p=0.99    0.9192
```

> **Result VI.2 (Landauer is per dropped coordinate). MEASURED / INTERPRETATION.**

> The erase cost of a site is its entropy relative to the reader, not a property of the bit or its value. A reader holding the section pays nothing, because for it the operation is a rotation, and rotations are free. The bill exists only in a frame that dropped a coordinate, and it is exactly the size of what was dropped, at the site it was dropped.

This closes the program's central omission. There is no ERASE instruction because erasure is never free and never local-to-the-bit; it is the relocation of a coordinate into degrees of freedom the reader stopped tracking, and the register's entropy loss equals the bath's entropy gain to the last digit. The register forgets; the universe does not. Nothing is erased — which is the change-mandate, C1, read at the finest scale: a valid continuation cannot destroy the history that produced the state, so the history goes into the fibre and is billed for at the boundary.

VI.4 The ontological claim, at the altitude the evidence supports

The framework from which this thesis grew is ambitious: it holds that reality is a self-consistent continuation field, that objects and laws and observers are emergent compressions of one transition structure, and that mathematics is universal because any system maintaining distinctions must implement equivalent continuation logic. This thesis does not prove that. It proves something narrower and offers the framework as the reading that the narrow results invite.

What is proved: that a continuous geometry can supply a family of possibilities while an arithmetic condition selects one member, demonstrated with full rigor in one case (Part IV); that the selecting operation appears, computed, in unrelated substrates (Part V); that the observable belonging to a linear operator read by a weight sequence is the reader's generating function at the operator's pole, which belongs to the pair and not to either alone (Part IV.1); and that the thermodynamic cost of computation is the cost of dropping a coordinate, not of manipulating a bit (Part VI.3).

What is therefore reasonable to say, and no more: if computation is the substrate-independent operation these results exhibit, then a physical system realizing a given operator will expose the invariant belonging to that operator, because the invariant belongs to the operator rather than the substrate. The constant a machine reads is a property of the pair (dynamics, reader), and π in particular is a level set that any such pair crosses. This is a conditional, and its antecedent — that a given physical system implements a given operator — is exactly what the thesis does not assert for any specific system.

> **The central conditional. INTERPRETATION.**
>
> The observed constant belongs to the pair (dynamics, reader), not to either in isolation. A physical realization of an operator reads the invariant of that operator because the invariant is the operator's, not the substrate's. Whether any particular physical system realizes any particular operator is a separate empirical question, untouched here.

> *The strongest sentence this thesis can stand behind is not that nature computes π, and not that hashing is biology. It is that a constraint defines an admissible evolution, a readout defines which invariant becomes visible, and the visible invariant belongs to the pair. Everything else in the framework is the reading this makes available — offered honestly as a reading, held separately from the theorems, and left open.*

VI.5 What remains open

The thesis is bounded by named gaps, and naming them precisely is part of its method. The general form of the void rank–nullity decomposition (Result II.3) is a conjecture, proved only per construction. The uniqueness theorem (Result IV.3) holds within one explicitly stated family and excludes nothing outside it; whether the period relation $\pi \in \mathbb{Q}\text{-span}\{I_a\}$ holds only at (8,16) across all bases is open and Baker-hard, and the companion period-relation paper treats it at length. The extension of the generating-function schema to nonlinear evolutions (Part V.4) is undeveloped and is stated as a debt rather than a claim. And the ontological conditional's antecedent is empirical and untested.

None of these gaps is hidden, and none is load-bearing for the results below it. That is the structure the framework insists on: a killed claim is a constraint pointing somewhere, the audit trail keeps the seventeen that were killed along the way, and the boundary of what is proved is drawn in the same ink as the proofs. The universe this thesis describes is incomplete in the one way that matters — not because something is missing from it, but because any reader of it is finite, and the gap is the reader's, and it moves when the reader moves while the object does not.

End of program. No halt instruction was reached, because the change-mandate does not provide one.